

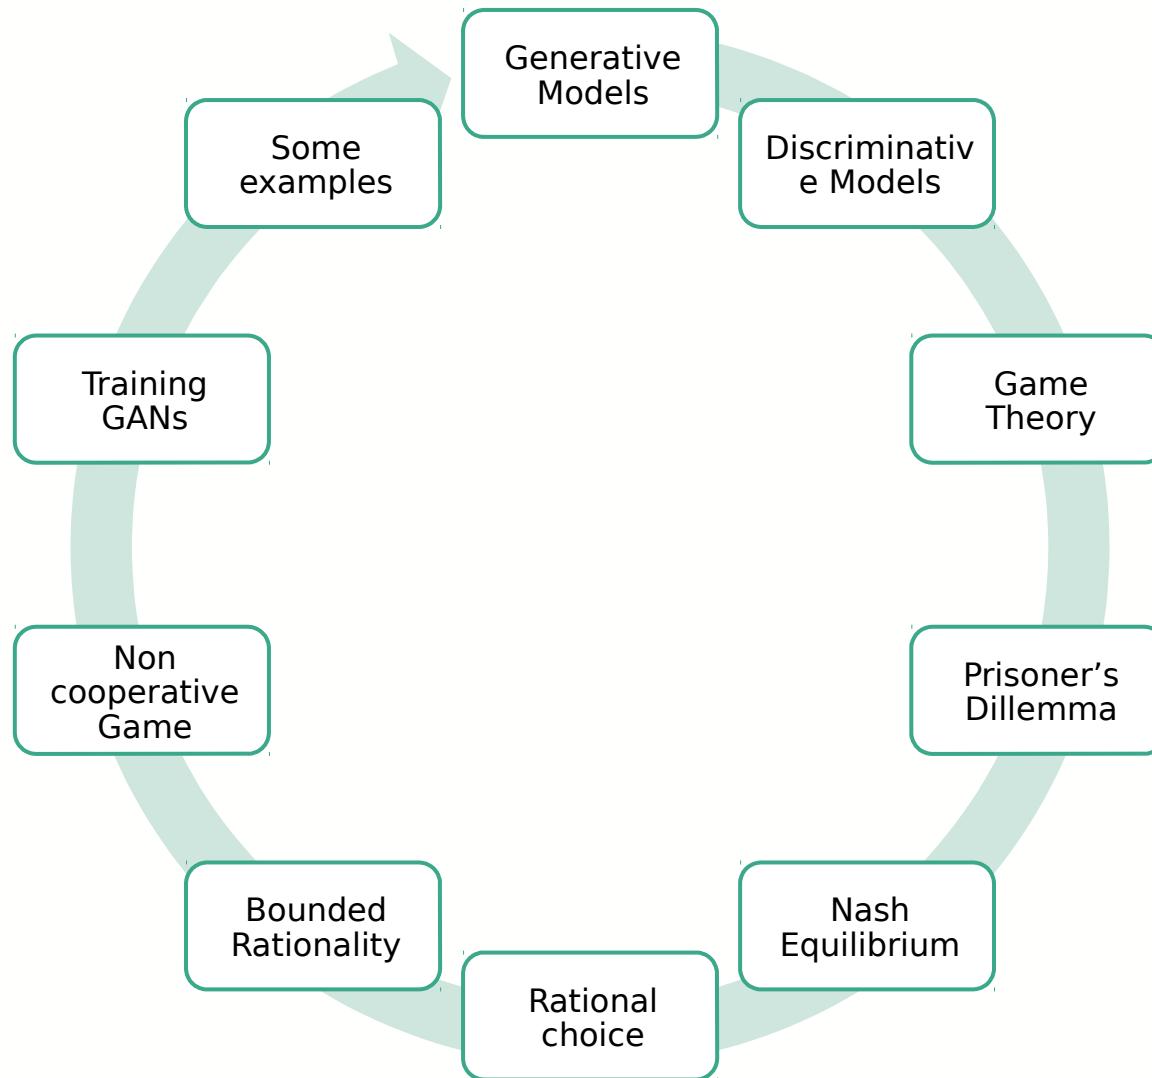
# Generative Adversarial Networks

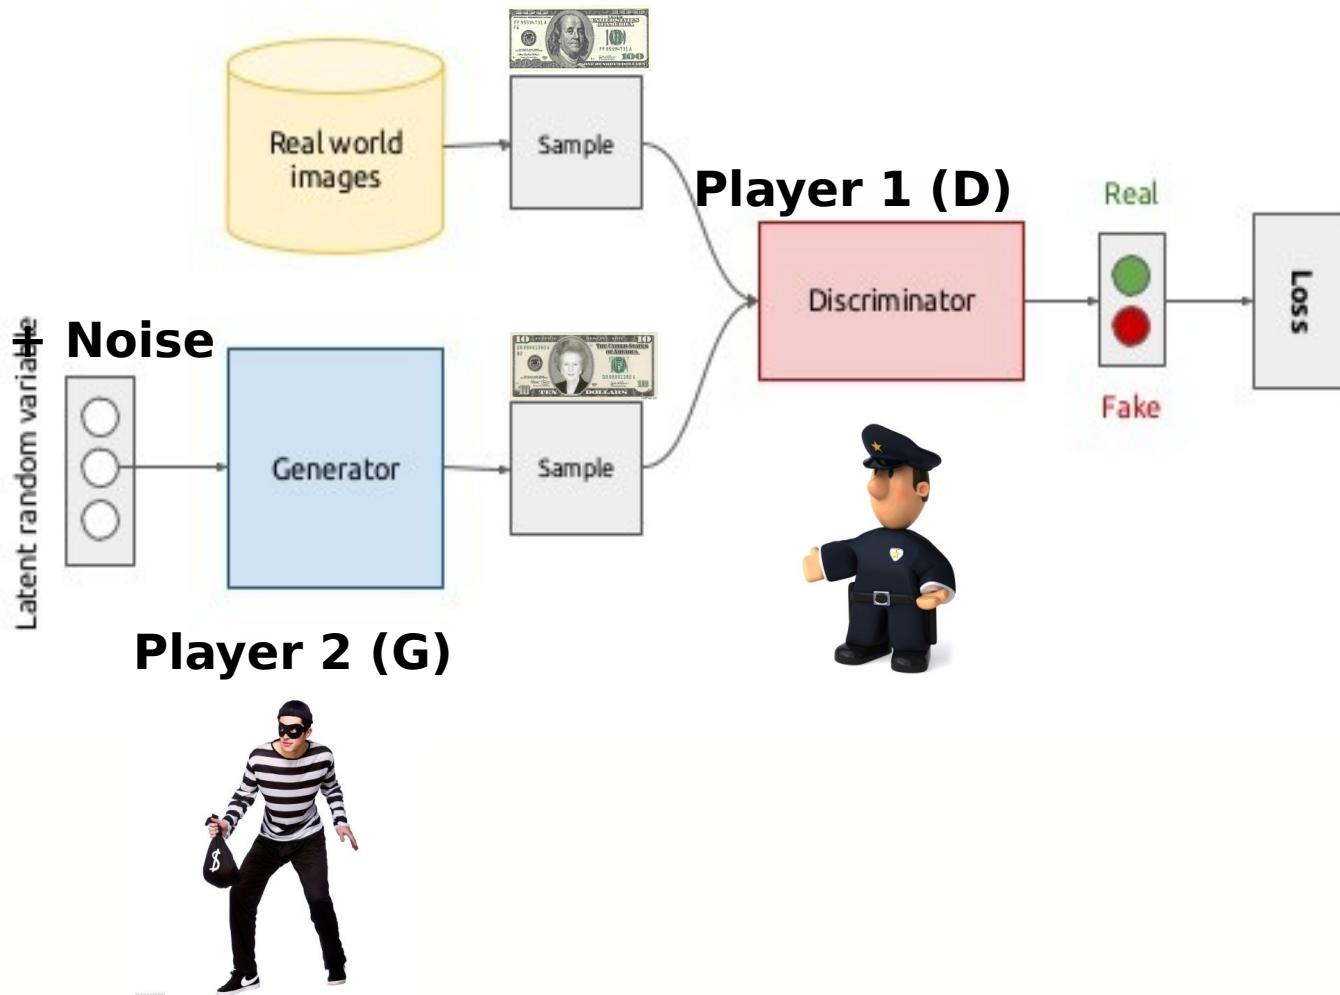
**Rubens Zimbres**

**Data Scientist**

**Machine Learning for IoT at Vecto Mobile**





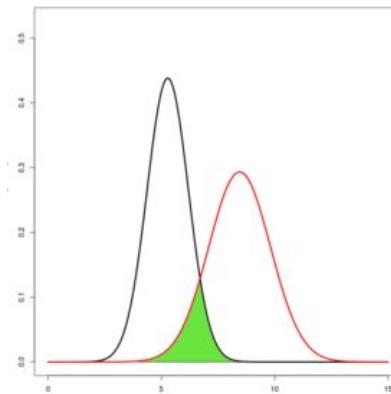
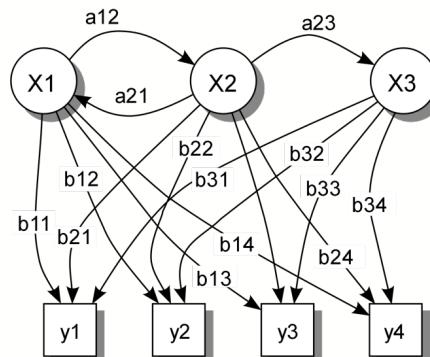


# Generative Models

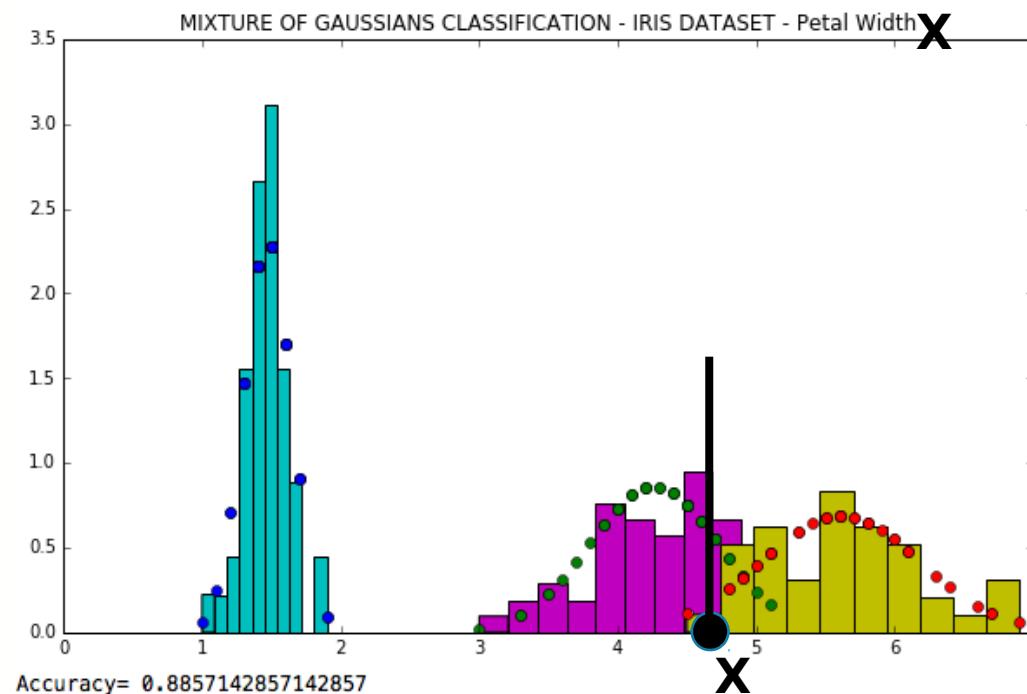
- Probabilistic, samples from joint distribution  $P(x,y)$
- Generate samples (latent variables) following a distribution
- Maximum Likelihood
- Range of values belong to a class

# Generative Models

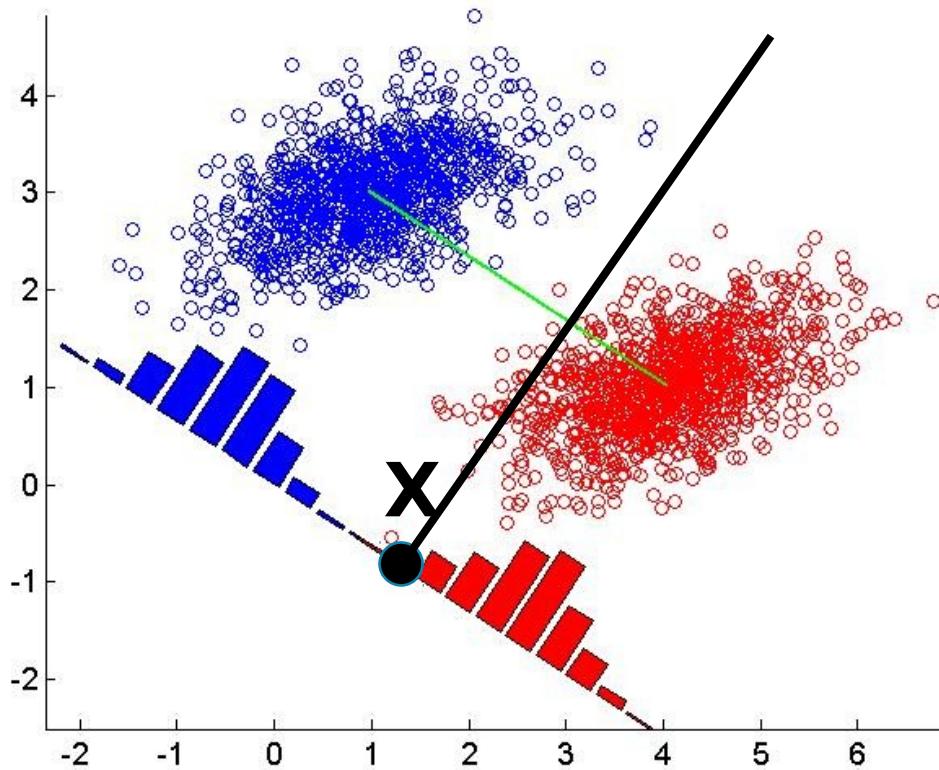
- Latent Dirichlet Allocation (NLP)
- Mixture Gaussians
- Hidden Markov Model
- VAE
- RBM
- GANs



# Mixture of Gaussians



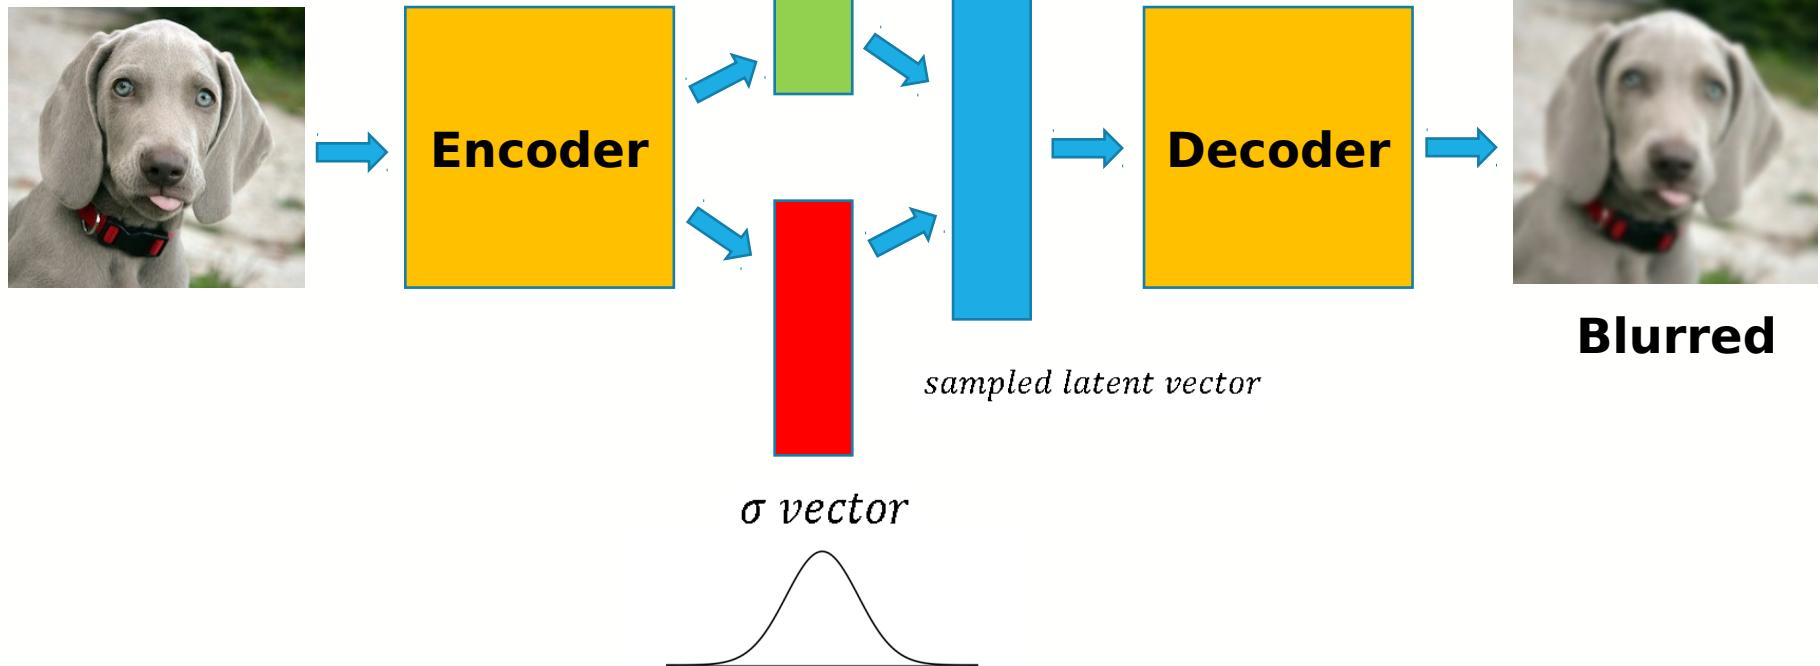
# Linear Discriminant Analysis



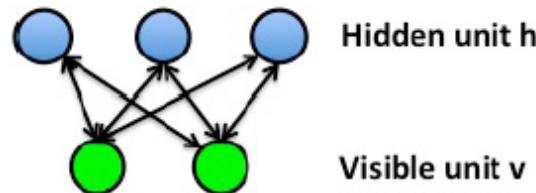
# Variational Autoencoders

ckprop

llback-Leibler divergence KL (surprise)



# Restricted Boltzmann Machines



$$E(v, h) = - \sum_{\text{Energy}} v_i h_j w_{ij}$$

Binary state      Binary state  
 of visible      of hidden

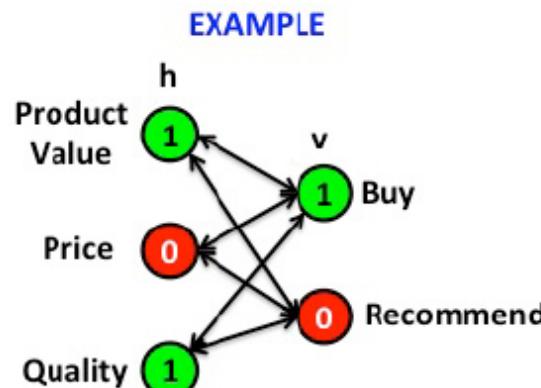
Probability of joint configuration

$$p(v, h) = \frac{e^{-E(v, h)}}{\sum e^{-E(u, g)}}$$

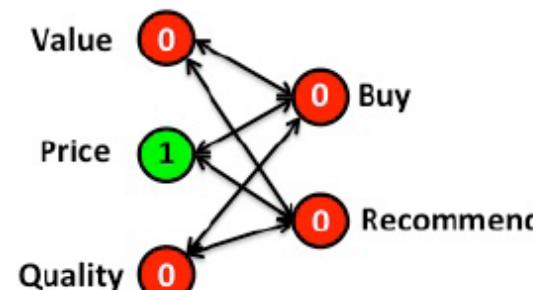
Energy of joint configuration  
 Energy of the system

## RBM Learning (update in parallel)

$$\Delta w_{ij} = \eta \cdot \frac{\partial}{\partial x} \log(p(v)) = \eta \cdot (< v_i h_j >^0 - < v_i h_j >^1)$$

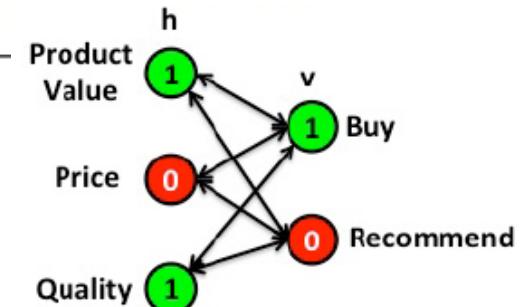
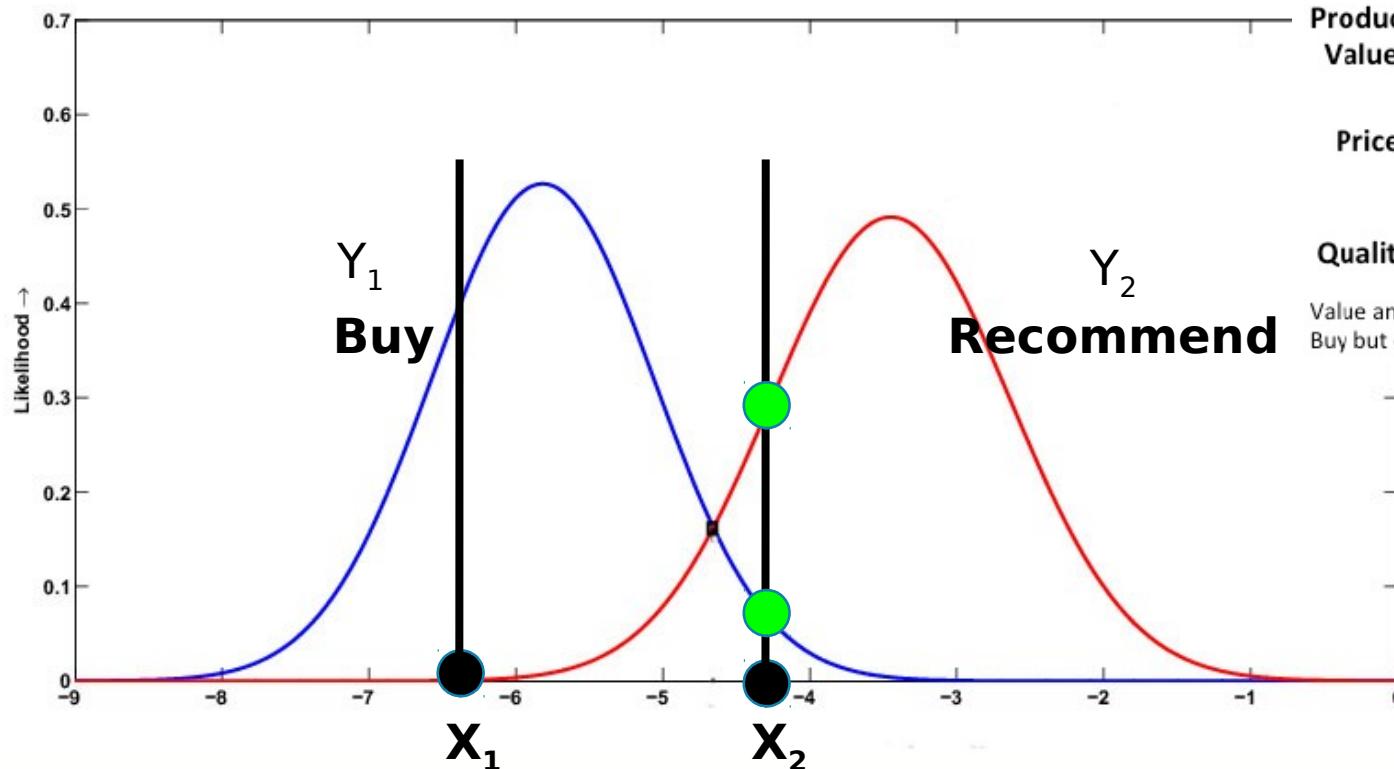


Value and Quality above average, but too expensive  
 Buy but does not Recommend to friends



Poor Value and Quality and affordable  
 Does not Buy and does not Recommend to friends

# Generative Models



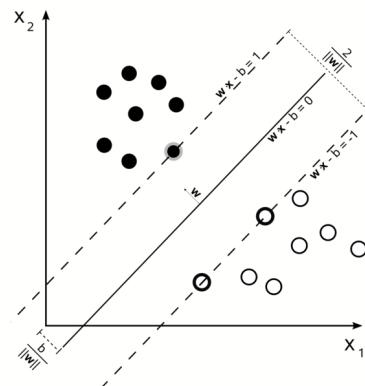
Value and Quality above average, but too expensive  
 Buy but does not Recommend to friends

# Discriminative Models

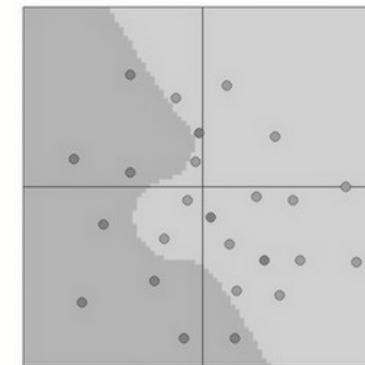
- Observed variables: target, infer outputs, conditional
- Function maps from  $x$  to  $Y$
- Conditional density function
- $P(Y|x)$  subsample
- Decision boundary
- Computationally expensive

# Discriminative Models

- Logistic Regression



- SVMs



- Neural Networks

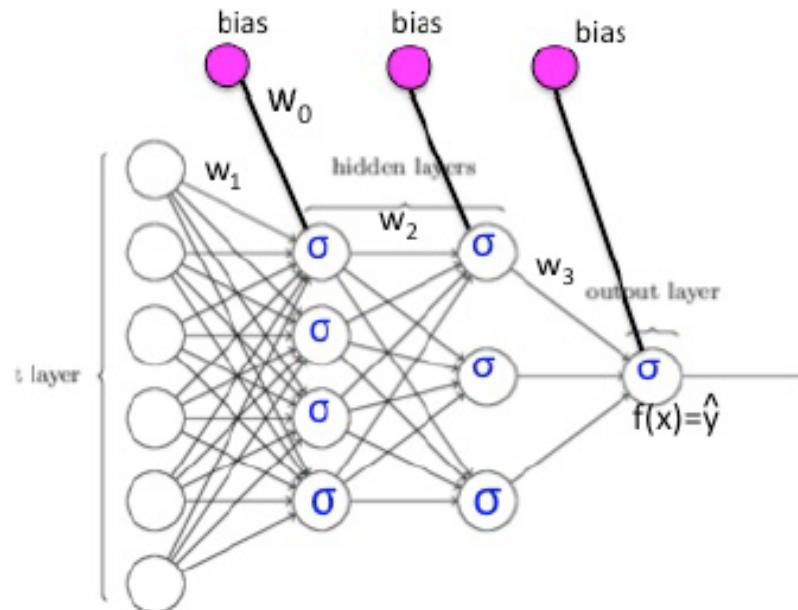
$$\text{Odds Ratio} = \log\left(\frac{P}{1-P}\right) = mx + b$$

$$\text{Odds} = \left(\frac{P}{1-P}\right) = e^{mx+b}$$

### Cost function Logistic Regression

$$J(\theta) = -\frac{\sum y \cdot \log(\hat{y}) + (1-y) \cdot \log(1-\hat{y})}{n}$$

$$\text{where } \hat{y} = \frac{1}{1 + e^{mx+b}}$$



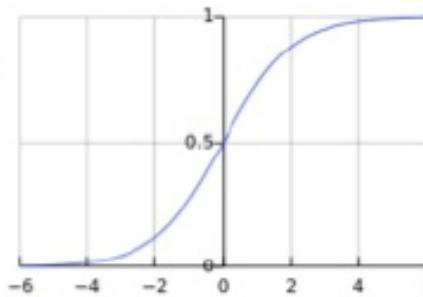
### Sigmoid Activation Function σ

### Cost function Neural Networks

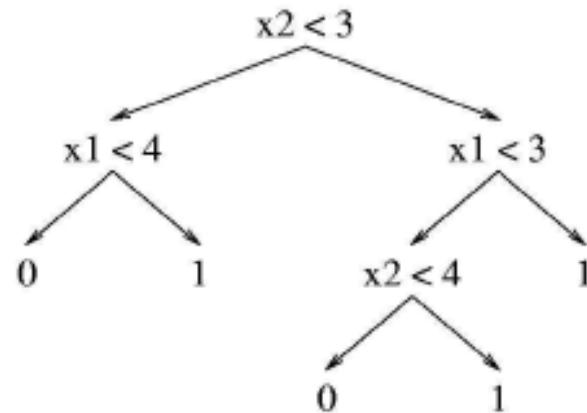
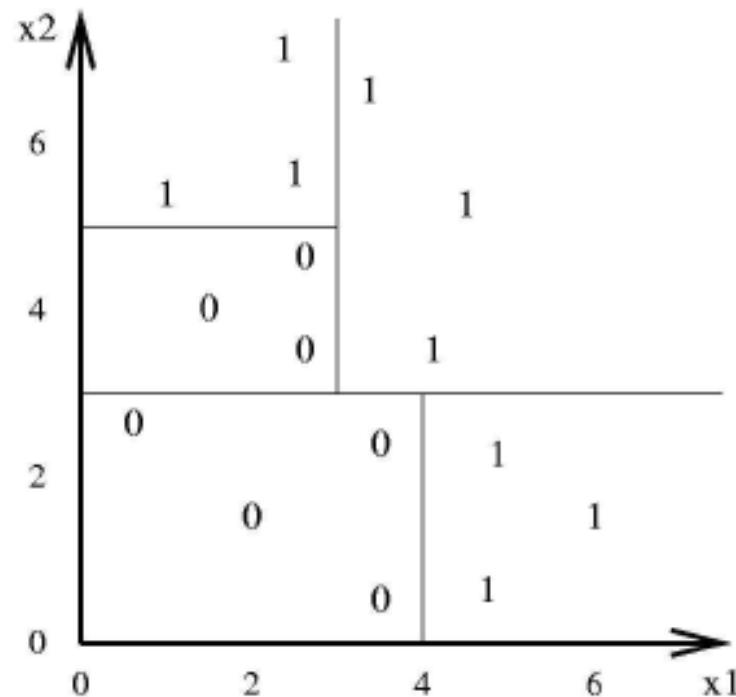
$$J_{\theta} = \frac{\sum_{i=1}^n \sum_{k=1}^k t_k \cdot \log(o) + (1-t) \cdot \log(1-o)}{N} + \frac{\lambda \sum_{l=1}^L \sum_{i=1}^I \sum_{j=1}^{j+1} \theta_{ji}^2}{2N}$$

L2 Regularization term

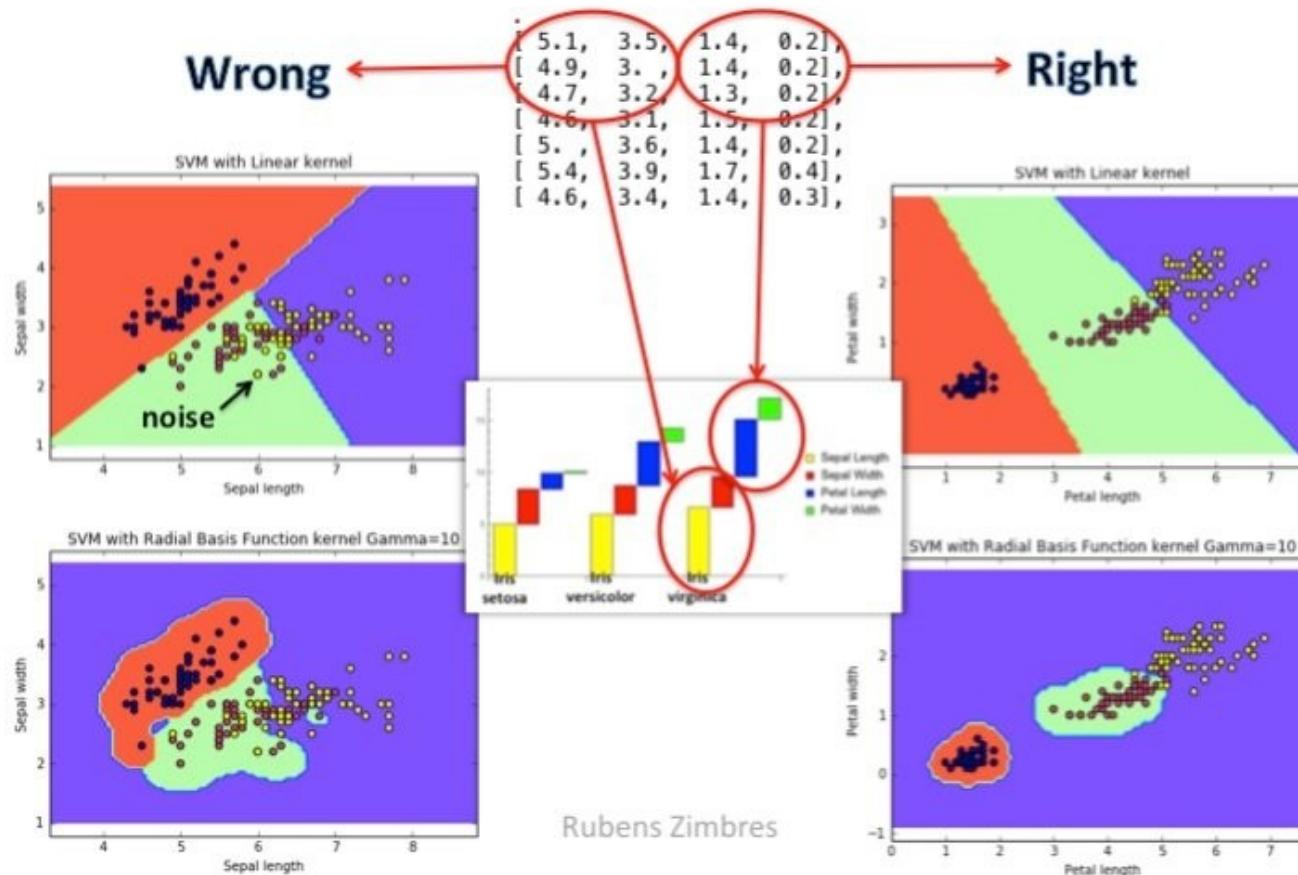
Where y = t = Target and y = o = Output



# Decision Trees



# Occam's razor Principle



# Generative vs Discriminative

- Generative learns joint probability distribution =  $P(x,y)$

$$\frac{\text{Probability}}{\sum \text{samples}}$$

- Discriminative learns conditional probability distribution =  $P(y|x)$

$$\frac{\text{Probability}}{\text{subsamples}}$$

# GAN Advantages over Generative Models

- D & G = Multi Layer Perceptron
- Regular backpropagation like VAE (G and D differentiable)
- Does not necessarily involve Maximum Likelihood estimation
- No need to use Markov Chains
- Subject but robust to overfitting

# Game Theory

- 2 Player Game - Zero Sum
- Minimax solution: Nash Equilibrium
- Unique solution:  $D=1/2$  everywhere

		PLAYER B	
		COOPERATE	DEFECT
		COOPERATE	A: 1 year jail B: 1 year jail
PLAYER A	COOPERATE	A: 10 years jail B: 0 years jail	
	DEFECT	A: 0 years jail B: 10 years jail	A: 5 years jail B: 5 years jail

### PAYOUT MATRIX

		Discriminator	
		Cooperate	Defeat
Generator	Cooperate	(1,1)	(0,1)
	Defeat	(1,0)	(0,0)

# Nash Equilibrium

- Sub optimal
- Non cooperative
- Emulate human behavior: bounded rationality (Herbert Simon)
- Simultaneous but not equivalent \*\* (GANs)

# GANs

- Semi Supervised Learning: missing labels
- Inverse Reinforcement Learning

## MARKOV DECISION PROCESS

$$U_s = R_s + \delta \max_a \sum_{s'} T(s, a, s'). U(s')$$

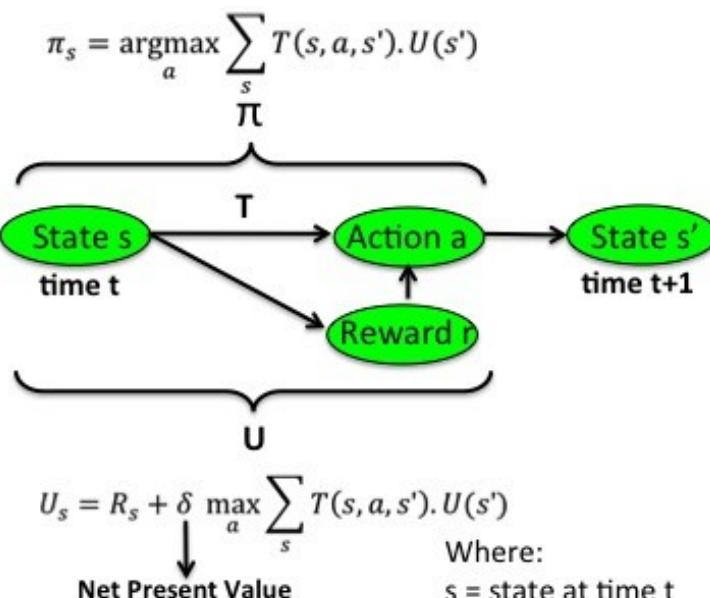
$$\pi_s = \operatorname{argmax}_a \sum_{s'} T(s, a, s'). U(s')$$

$$Q_{s,a} = R_s + \delta \max_{s'} \sum_s T(s, a, s'). \max_{a'} Q(s', a')$$

$$\hat{Q}_{s,a} \leftarrow \eta R_s + \delta \max_a Q(s', a')$$

# Reinforcement Learning

## Markov Decision Process



Prospect Theory  
 Expected Utility Theory  
 Rational Choice Theory  
 Game Theory

## Game Theory

### Prisoner's Dilemma

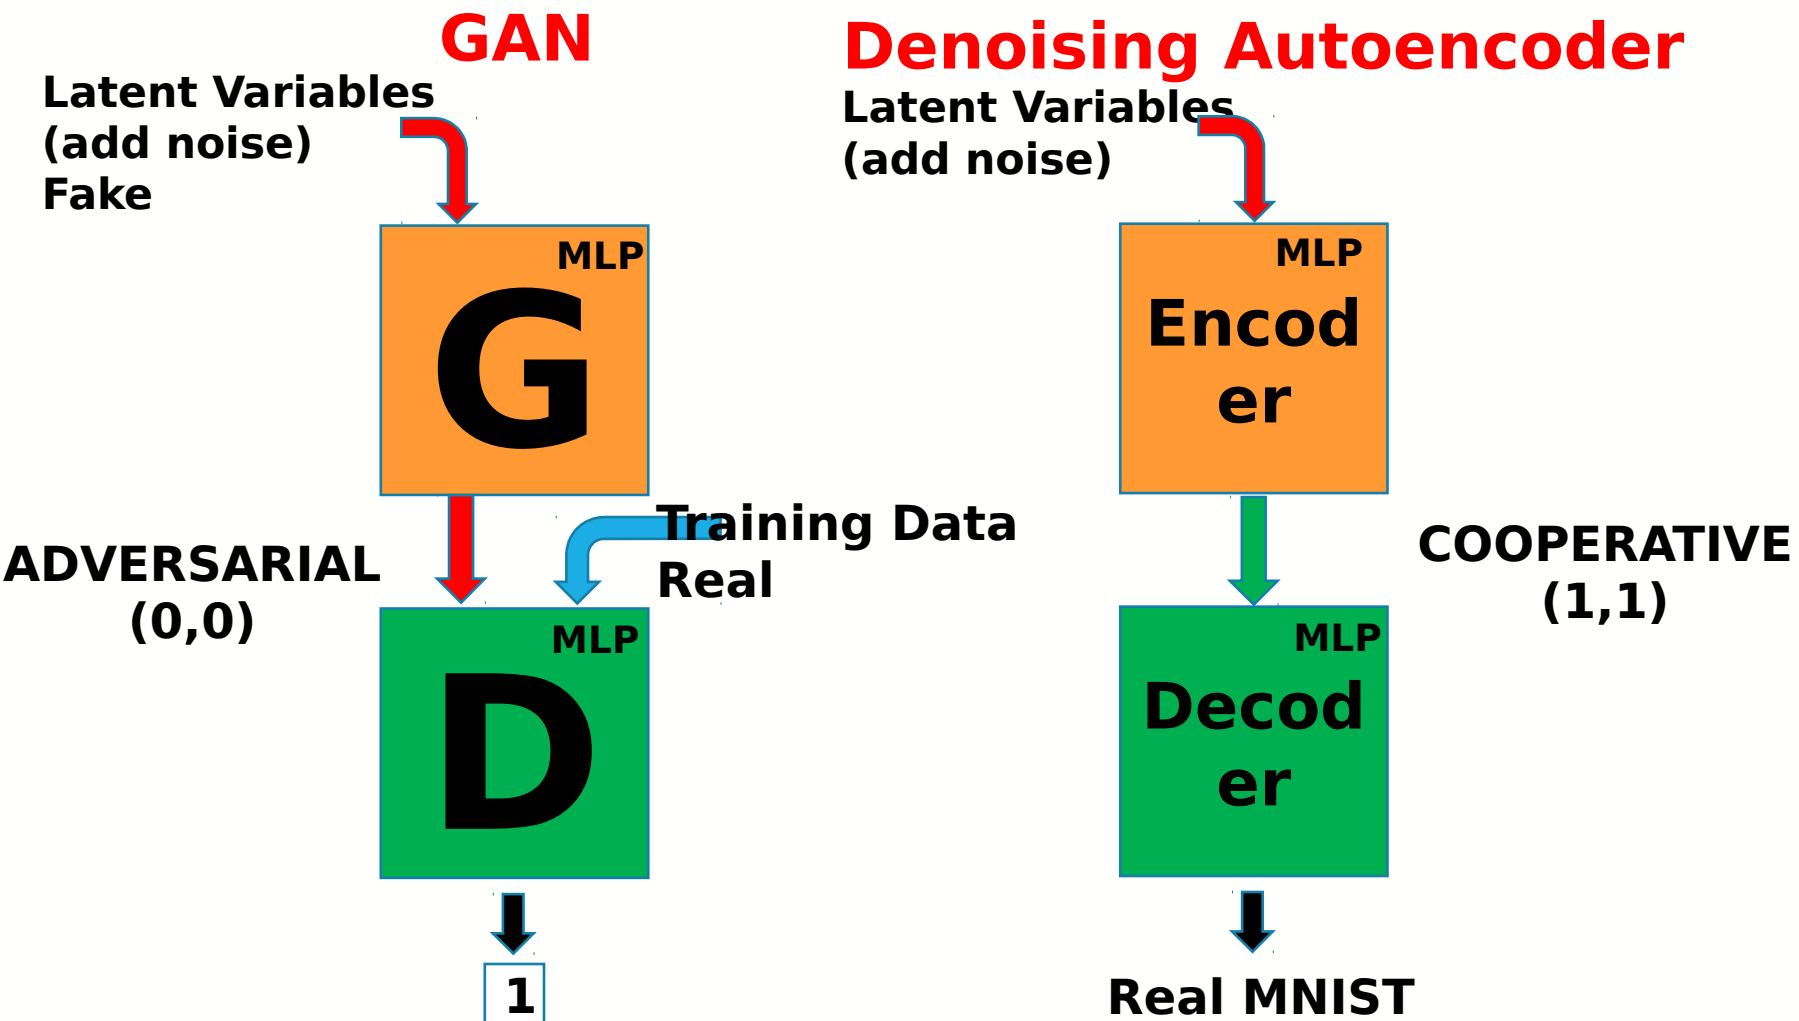
		B	
		Cooperate	Defeat
A	Cooperate	(-1, -1)	(-8, 0)
	Defeat	(0, -8)	(-5, -5)

### Static: Nash Equilibrium:

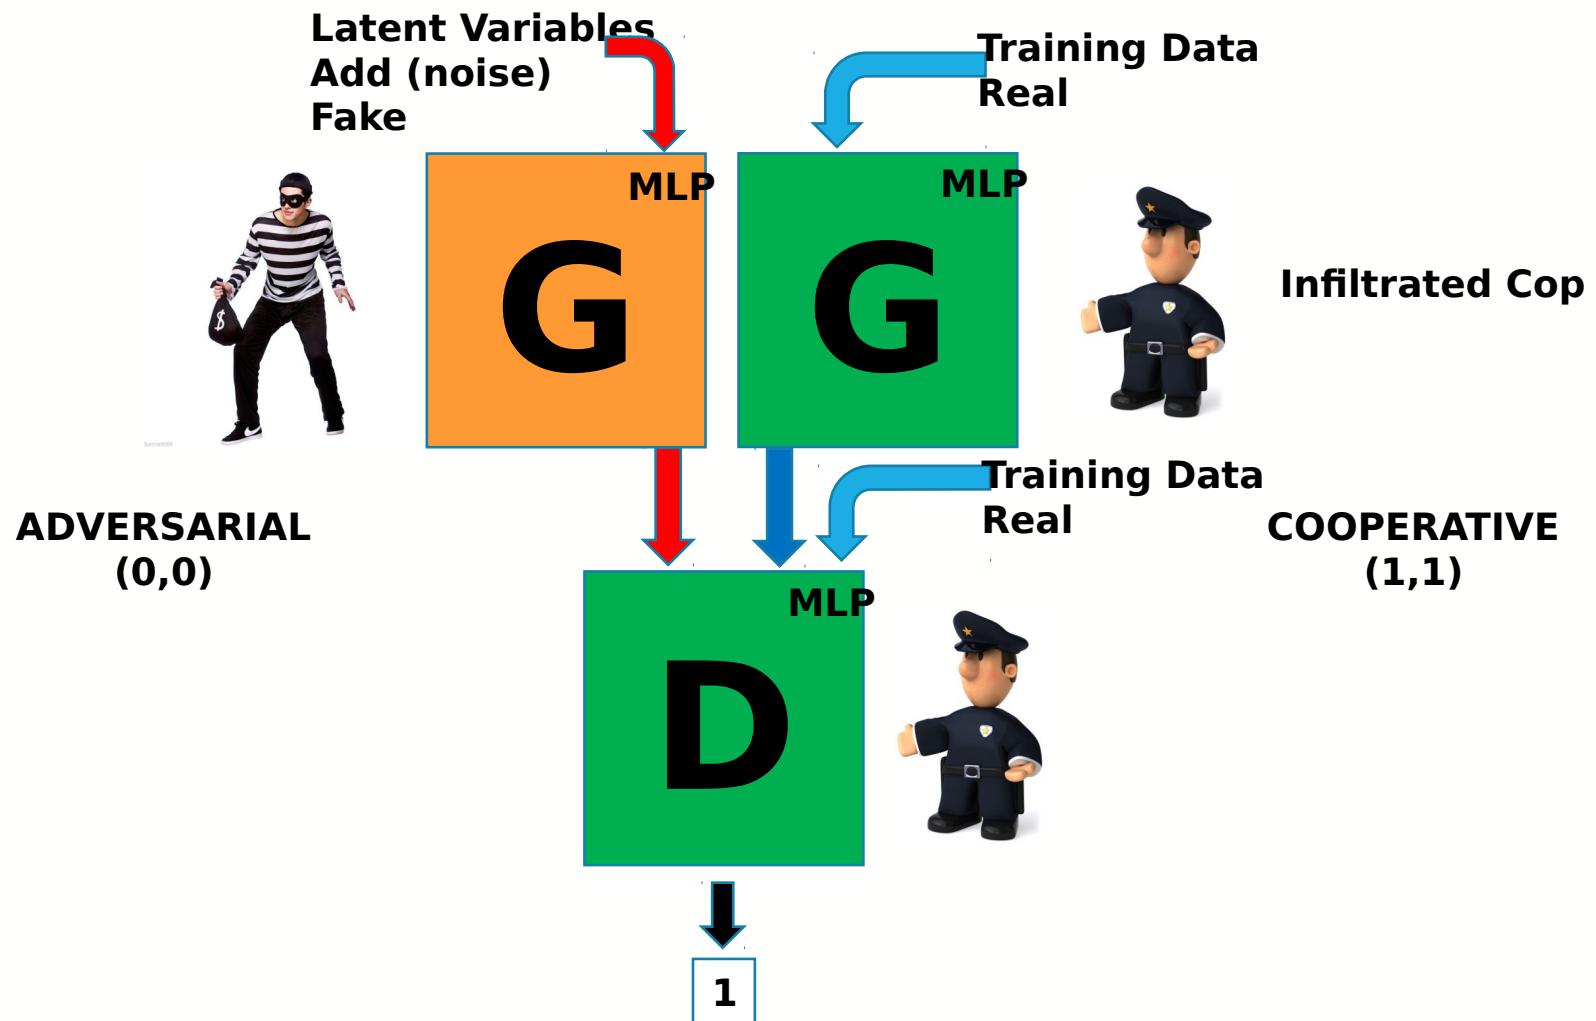
- “Best” strategy chosen by all players at same time
- No reason to change strategy
- Inefficient: (-1,-1) is better
- Herd behavior

### Dynamic: Iterative

- Pavlovian conditioning
- Tit-for-tat (Cold War times)
- Mutual Assured Destruction
- Mixed strategy (probability)



## Siamese GAN



# GAN Tuning

- Neural Network architecture
- Hyperparameters
- Game Theory strategies

# Neural Networks

## 24 Adjustements

### ARCHITECTURE

- Variables type
- Variable scaling
- Cost function
- Neural Network type:
  - RBM,FFN,CNN,RNN...
- Number of layers
- Number of hidden Layers
- Number of nodes
- Type of layers:
  - LSTM, Dense, Highway
  - Convolutional, Pooling...
- Type of weight initialization
- Type of activation function
  - Linear, sigmoid, relu...
- Dropout rate (or not)
- Threshold

### HYPERPARAMETER TUNING

- Type of optimizer
- Learning rate (fixed or not)
- Regularization rate (or not)
- Regularization type: L1, L2, ElasticNet
- Type of search for local minima:
  - Gradient descent, simulated annealing, evolutionary...
- Batch size
- Nesterov momentum (or not)
- Decay rate (or not)
- Momentum (fixed or not)
- Type of fitness measurement:
  - MSE, accuracy, MAE, cross-entropy,
  - precision, recall
- Epochs
- Stop criteria

Rubens Zimbres

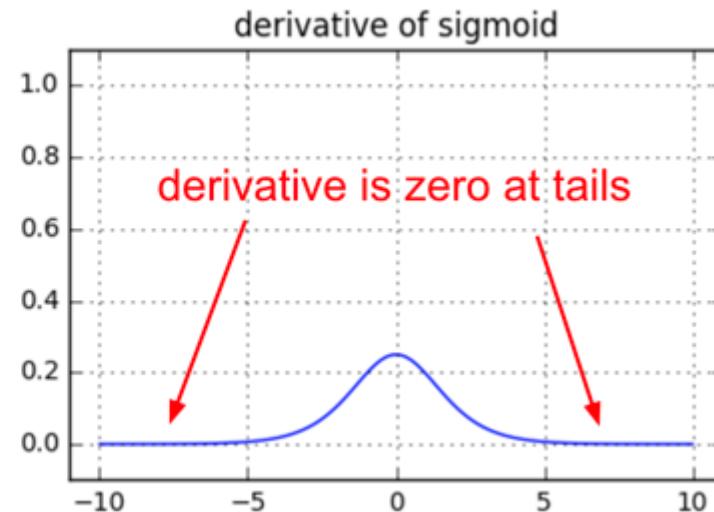
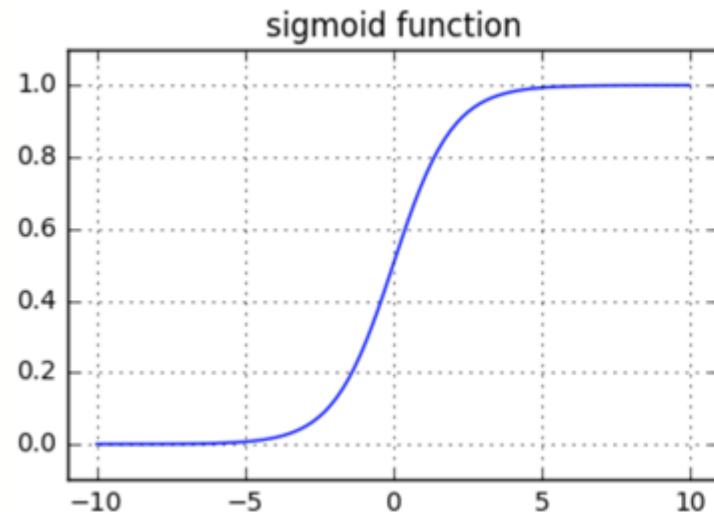
# GANs

- Train Simultaneously (no freeze)
- Train Discriminator one step
- Train Generator k steps (no Nash)
- Until  $p_G = p_{\text{data}}$  (global optimum)  $D(x) = P(x \in \text{data}) = 1/2$
- Discriminator cannot distinguish both distributions

# GANs

- $x$  comes from  $p_{\text{data}}$   $y=1$  (Real)
- $x$  comes from  $p_G$   $y=0$  (Fake)
- Optimize  $\theta$  rather than  $p_G$
- G: Relu (vanishing gradient)
- Dropout
- Supervised Learning in D:
- Beginning: high confidence
- Usually D is bigger/deeper than G

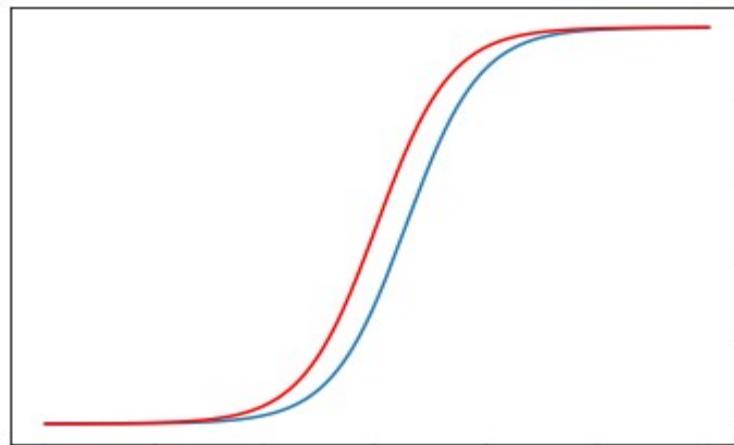
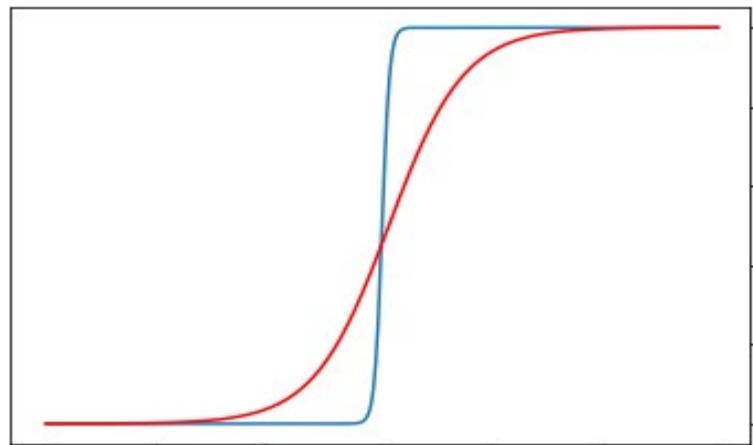
# Vanishing Gradient



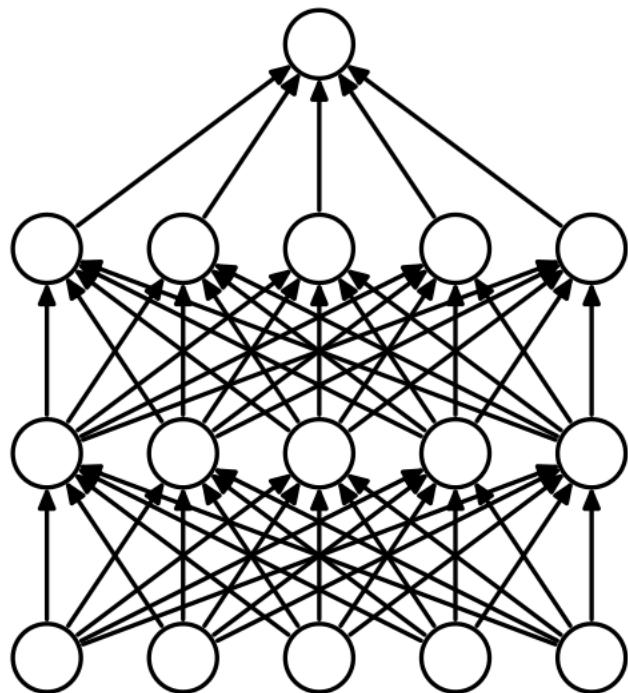
**-25% each activation**

Image source: Andrej Karpathy

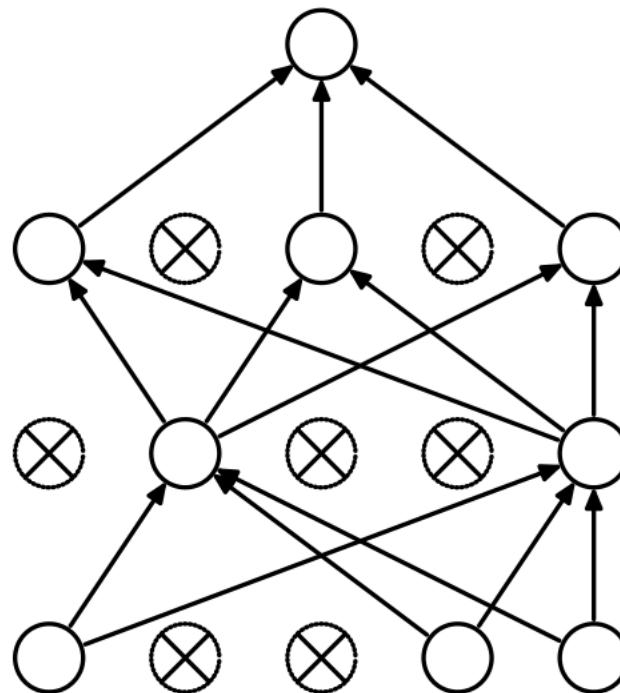
# Coefficient and Intercept: Regularization



# Dropout



(a) Standard Neural Net



(b) After applying dropout.

# Training GANs

$$V(G, D) = \mathbb{E}_{p_{\text{data}}} \log D(\mathbf{x}) + \mathbb{E}_{p_{\text{generator}}} (\log (1 - D(\mathbf{x})))$$

- Generator: Gradient Descent on V
- Discriminator: Gradient Ascent on V
- Minibatch
- Batch Normalization in G \*\*\*
- To minimize Cross Entropy
- Minimax Game
- Regular cross entropy

# Training GANs

- Discriminator minimize  $J^{(D)}(\theta_D, \theta_G)$  controlling  $\theta_D$
- Generator minimize  $J^{(G)}(\theta_D, \theta_G)$  controlling  $\theta_G$
- Cannot control each other

*D tries  $D(G(z)) \sim 0$*

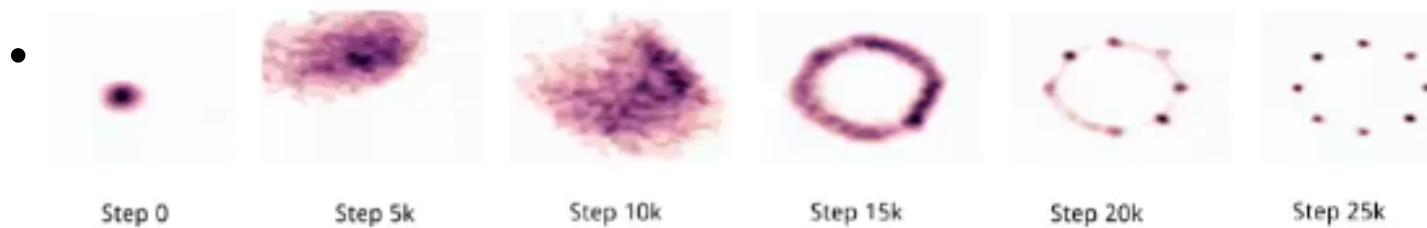
*G tries  $D(G(z)) \sim 1$*

$$J_D = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log(1 - D(G(z)))$$

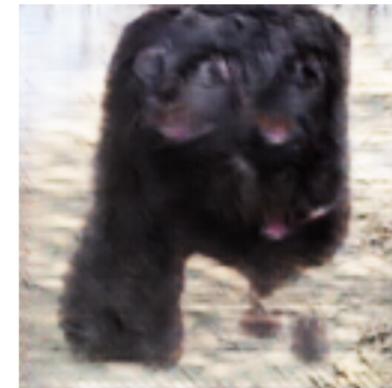
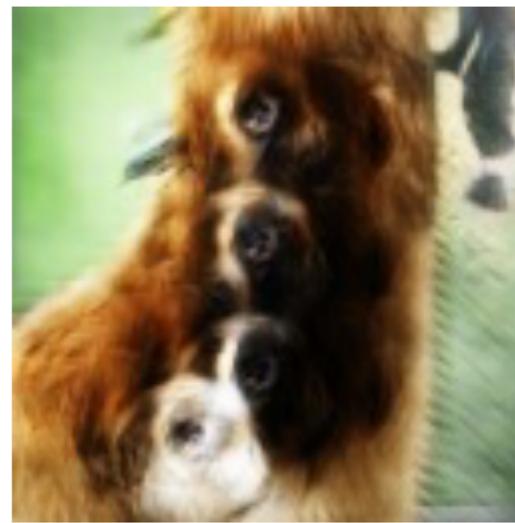
$$J_G = -J_D \quad J_G = -\frac{1}{2} \mathbb{E}_z \log D(G(z)) \quad (\text{G still learns})$$

# Training Challenges

- Non convergence: Nash Equilibrium, harder to optimize than objective function
- Mode collapse: G fails to output diversity (few good samples)
  - D converges to right distribution
  - G generates samples in the most probable point



# Training Challenges



# Tips and Tricks for GANs

- Normalize data (-1,1)
- Activation function Tanh output of Generative
- Sample from gaussian distribution instead of uniform distribution
- Develop different mini-batches in Discriminator:
  - All real
  - All fake
- Use Batch Normalization

# Tips and Tricks for GANs

- Use DCGANs or VAE+GAN
- Adam in Generator (exp. momentum decay)

(Adaptive Moment Estimation)

- SGD in Discriminator

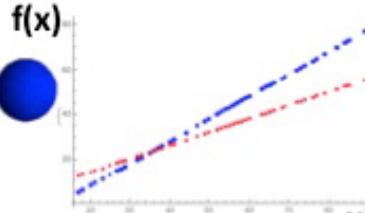
$$J(\theta_j) := \theta_j \pm \eta \cdot (\hat{y} - y)^2 \cdot x$$

- If Generator loss high = garbage to Discriminator
- Dropouts in Generator (overfitting)

# Gradient Descent for Linear Regression

$$f(x)_{t0} = \sum_{i=1}^n m_{i(t0)}x_i + b$$

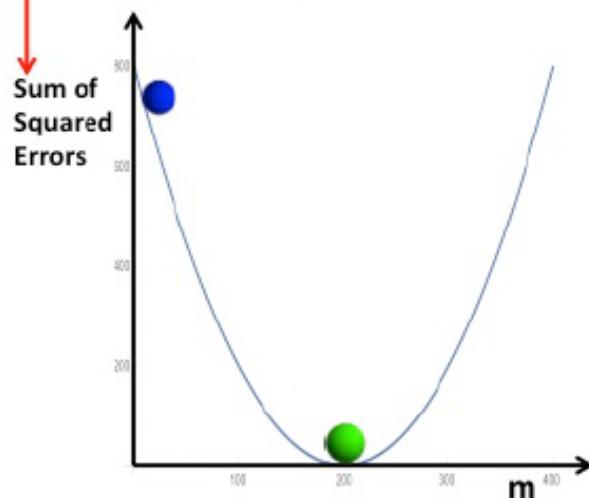
Linear Regression



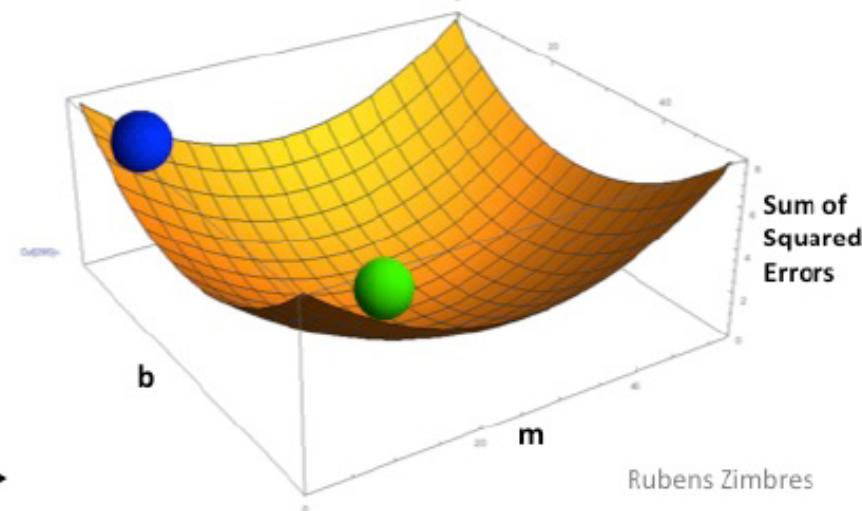
$$m_t := m_{t-1} \pm \eta \cdot \frac{\partial}{\partial x} \frac{\sum_{i=1}^n (m_{i(t0)}x_i - m_{i(t-1)}x_i)^2}{2n}$$

$$m_t := m_{t-1} \pm \eta \cdot \frac{\sum_{i=1}^n (m_{i(t0)}x_i - m_{i(t-1)}x_i)}{n}$$

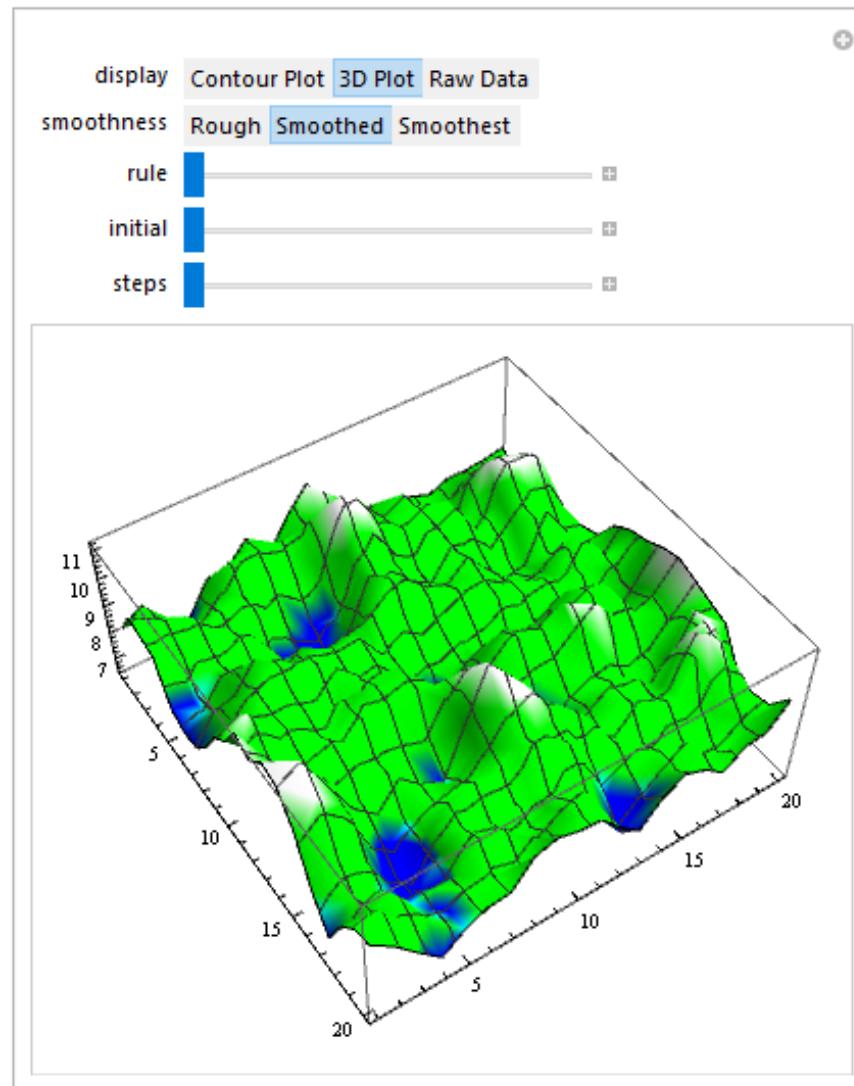
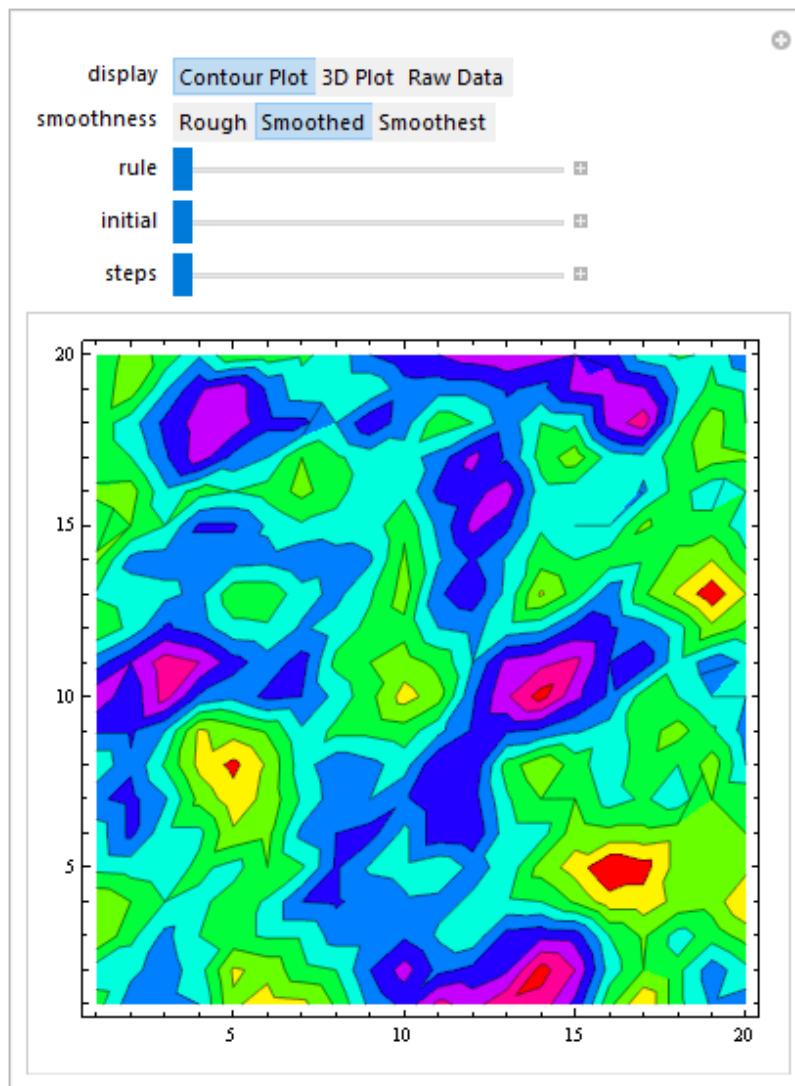
$\eta = \text{Learning Rate} \approx 0.01$



```
Plot3D[Table[{x^2 + y^2}, {x, -2, 2, 4/51}, {y, -2, 2, 4/51}], ImageSize -> 700]
```



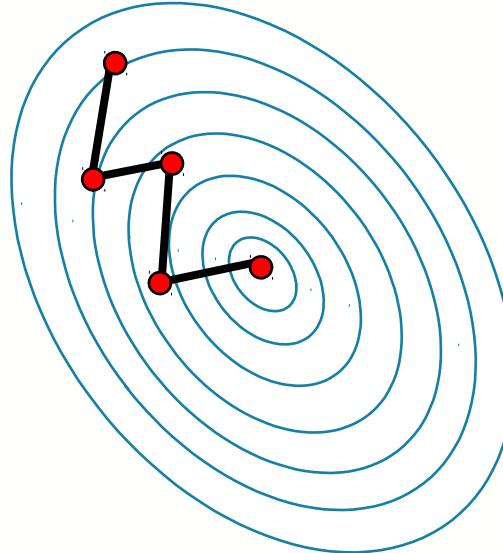
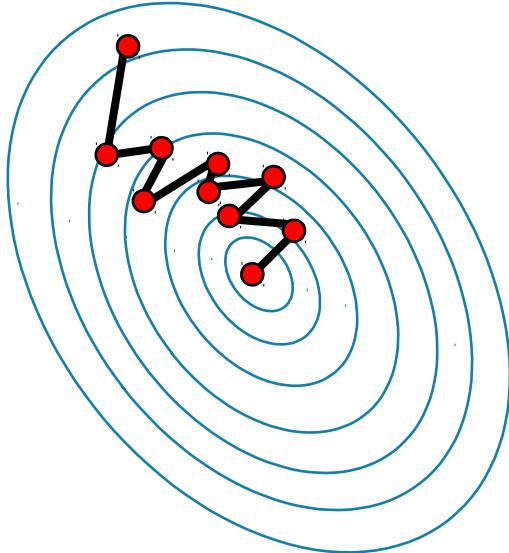
# DEEP LEARNING BRASIL SUMMER SCHOOL



# Gradient Descent

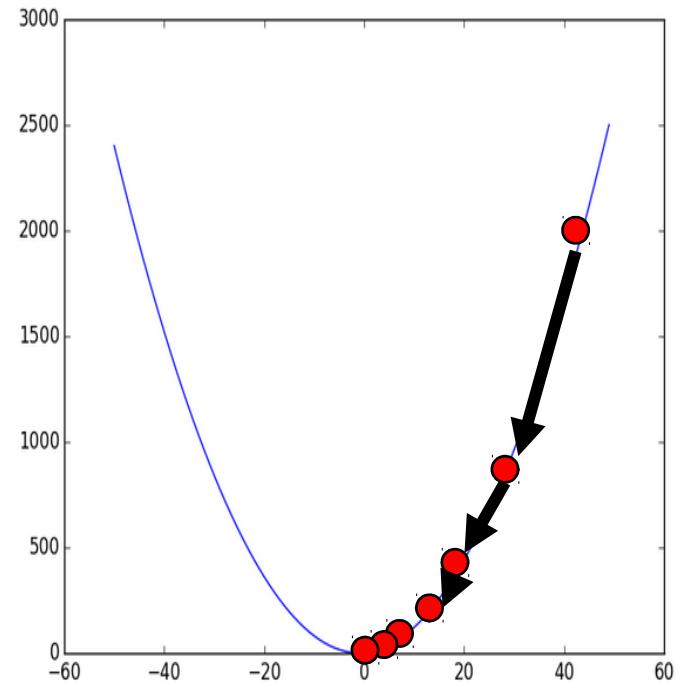
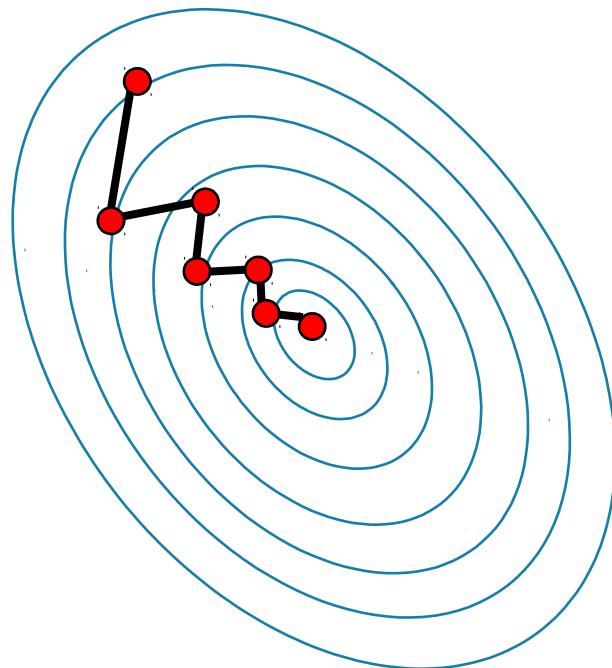
## Nesterov Momentum

$$\theta = \theta - (\gamma v_{t-1} + \eta \cdot \nabla J(\theta - \gamma v_{t-1}))$$

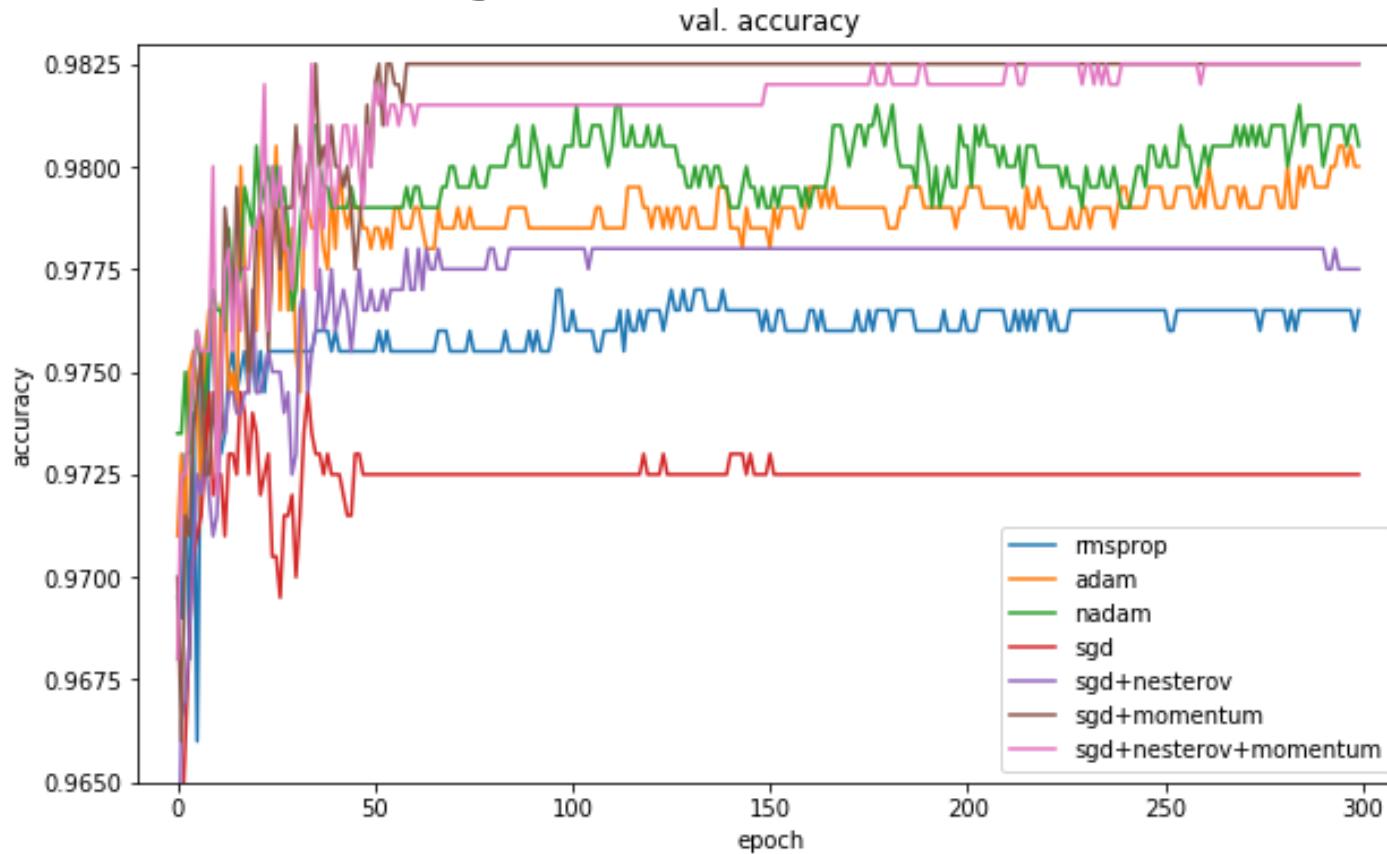


# Adam Optimizer

**Momentum w/ exponential decay according to derivative of error**

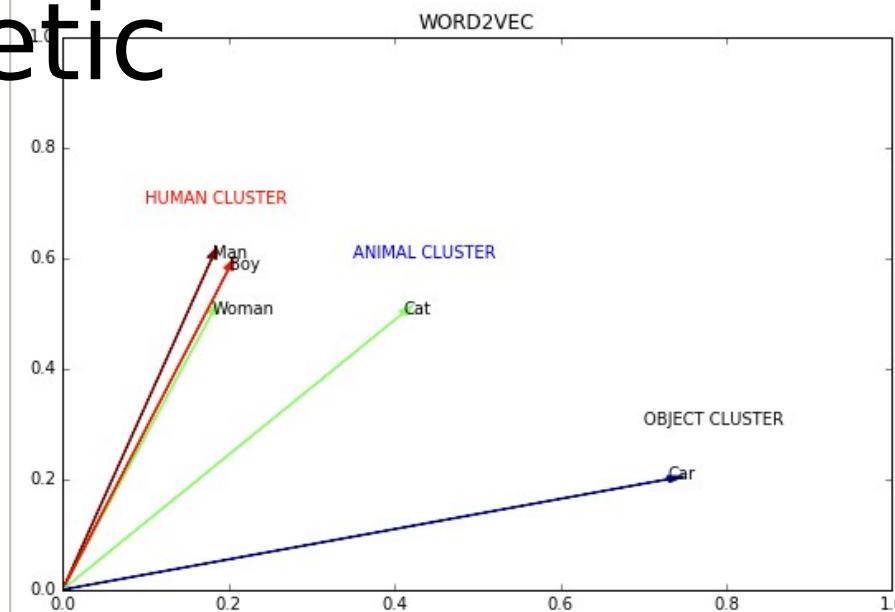


# Optimizers: generalization



Wilson et al, 2017

# GAN Examples: Vector Arithmetic



KNN Similarity with Man

Woman : 0.9

Man : 1.0

Boy : 0.971715728753

Cat : 0.74920127592

Car : 0.319926474563

Cosine Similarity with Man: Woman 0.882122164357

Cosine Similarity with Man: Man 1.0

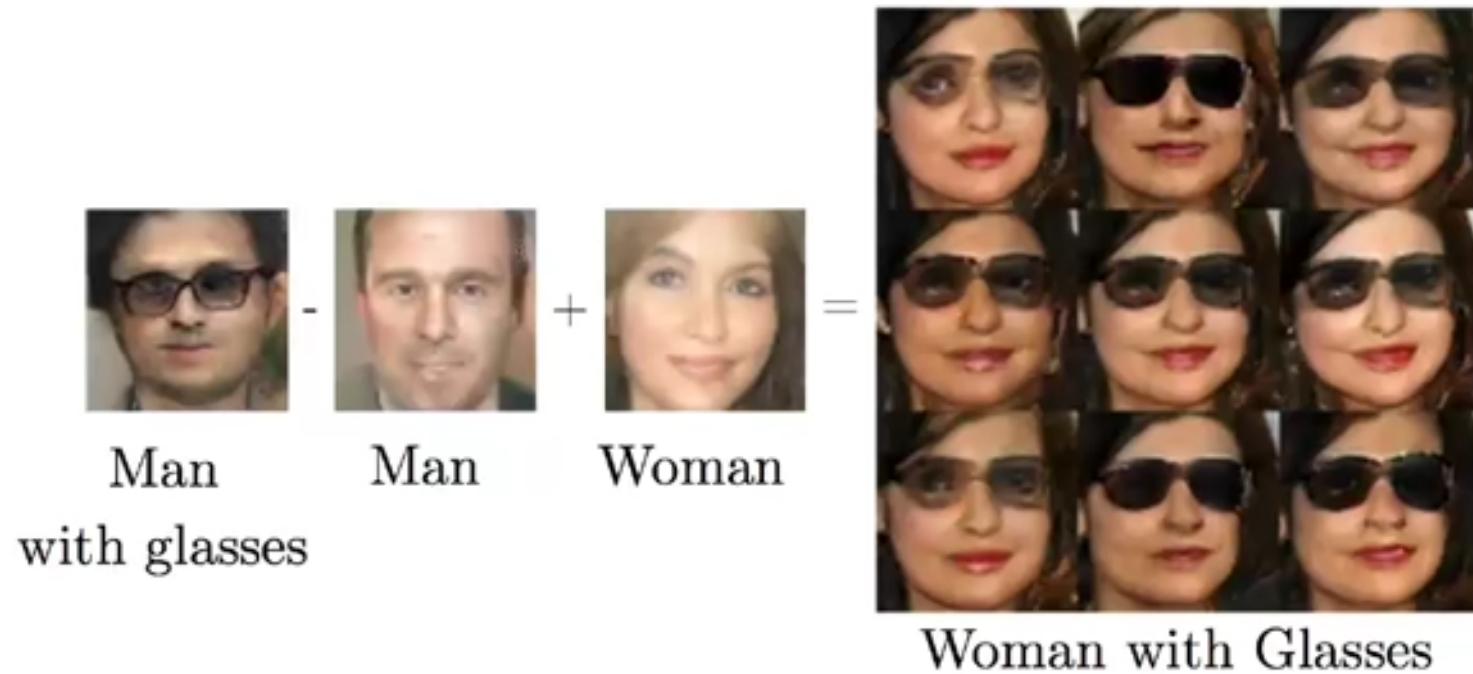
Cosine Similarity with Man: Boy 0.971120831449

Cosine Similarity with Man: Cat 0.757031631033

Cosine Similarity with Man: Car 0.437165180814

In [3]:

# Vector Arithmetic



Radford et al, 2015

# Matrix Multiplication

Color Guided Matrix Multiplication for a Binary Classification Task  
with  $N = 4$

Input Layer

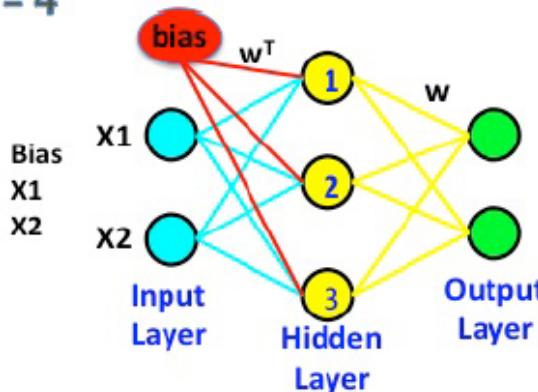
bias	X1	X2
1	0	1
1	0	0
1	0	0
1	1	0

$$\begin{bmatrix} 1 & 0 & 1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} .5 & .5 & .5 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \end{bmatrix} =$$

Weights  $w^T$  (transposed)

4 x 3      3 x 3

Go to Hidden Nodes



Hidden Layer

Bias	1	1	1
Node 1	.5	.5	.5
Node 2	.5	.5	.5
Node 3	1	1	1

$$= \begin{bmatrix} 1 & 1 & 1 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \\ 1 & 1 & 1 \end{bmatrix} \cdot \frac{1}{1 + e^{-(wx+b)}}$$

Sigmoid Function

Weights

.2	.1
.4	.1

$$= \begin{bmatrix} 1 & .3 \\ .5 & .15 \\ .5 & .15 \\ 1 & .3 \end{bmatrix} = \frac{1}{1 + e^{-(wx+b)}}$$

Output Layer

Sigmoid Function

$$= \frac{1}{1 + e^{-(wx+b)}}$$

Output

1	0
1	0
1	0
1	0

Rubens Zimbres

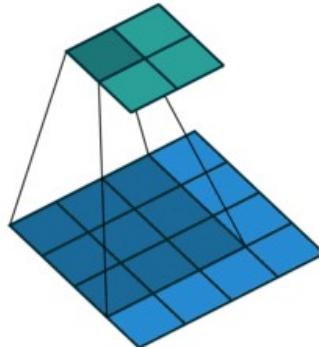
# DCGAN

- Deep Convolutional GAN
- Batch Norm not in end of Generator and NOT in Discriminator (x covariate shift)
- To increase dimension, Upsampling: `conv2D.T` with `stride > 1`
- Downsampling: average pooling + `conv2D` + `stride`

# Convolutions

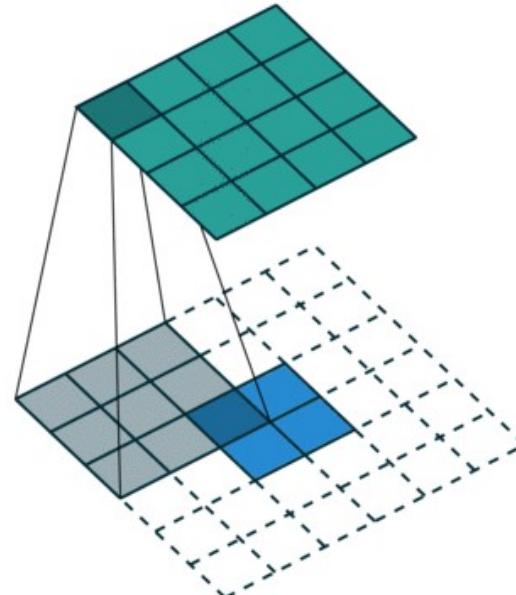
Convolution2D

3x3 kernel 4x4 input stride=1

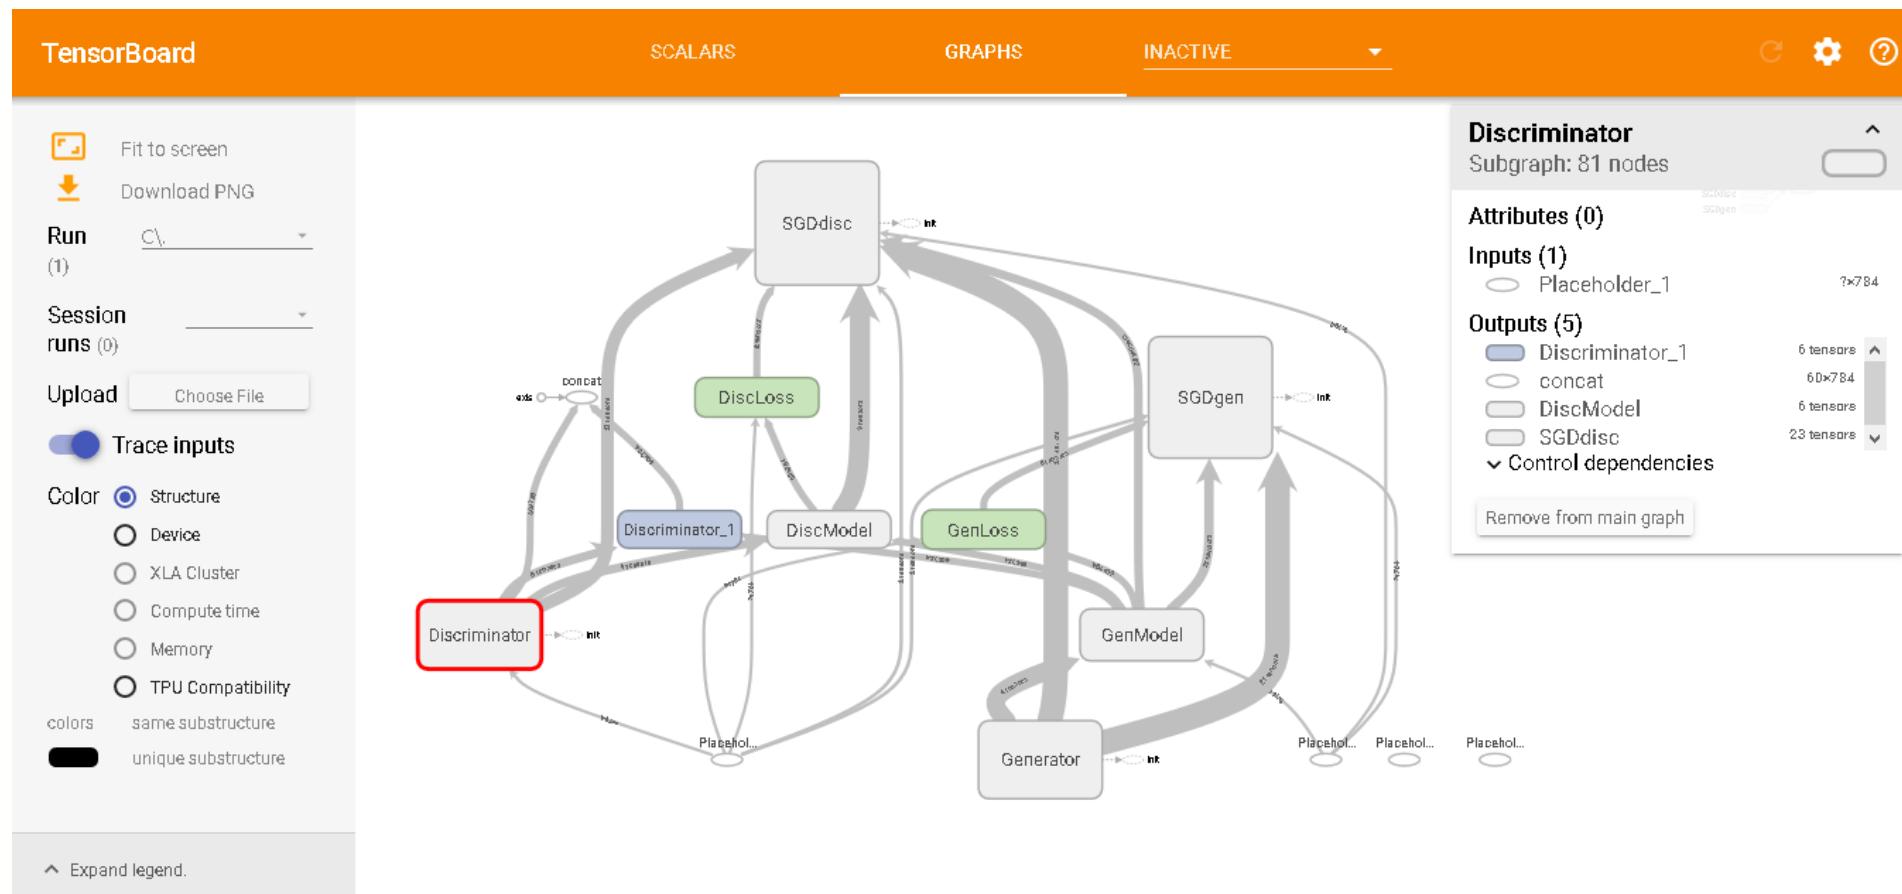


Convo2DTranspose

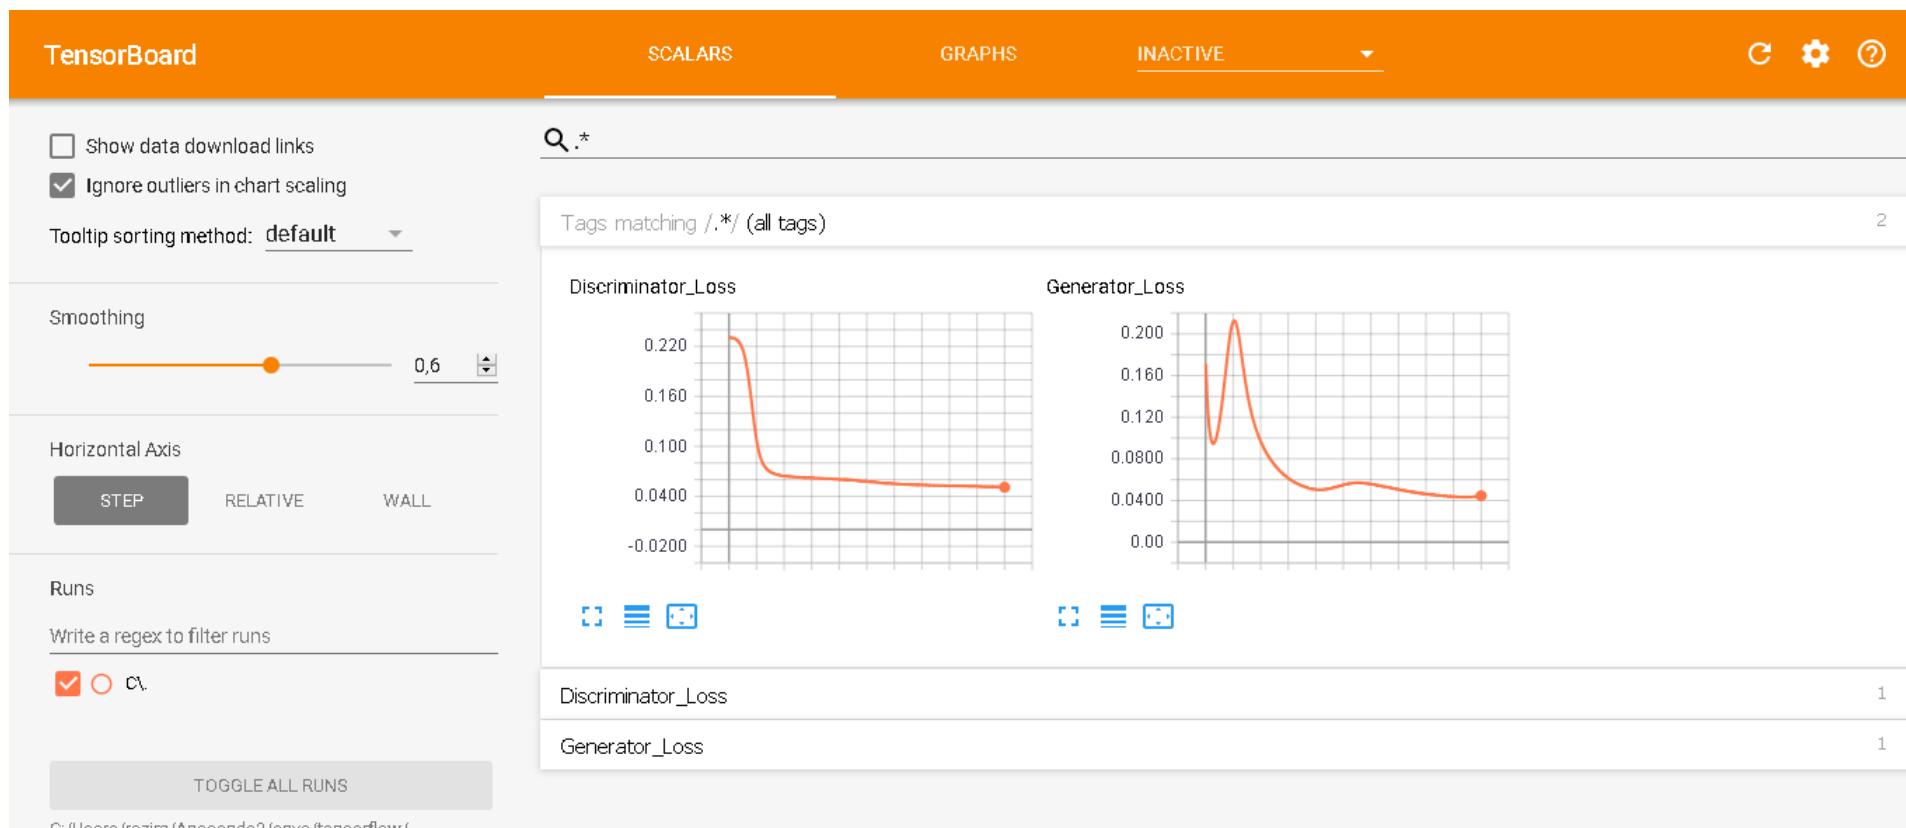
3x3 kernel 2x2 input stride=1



# DCGAN Architecture

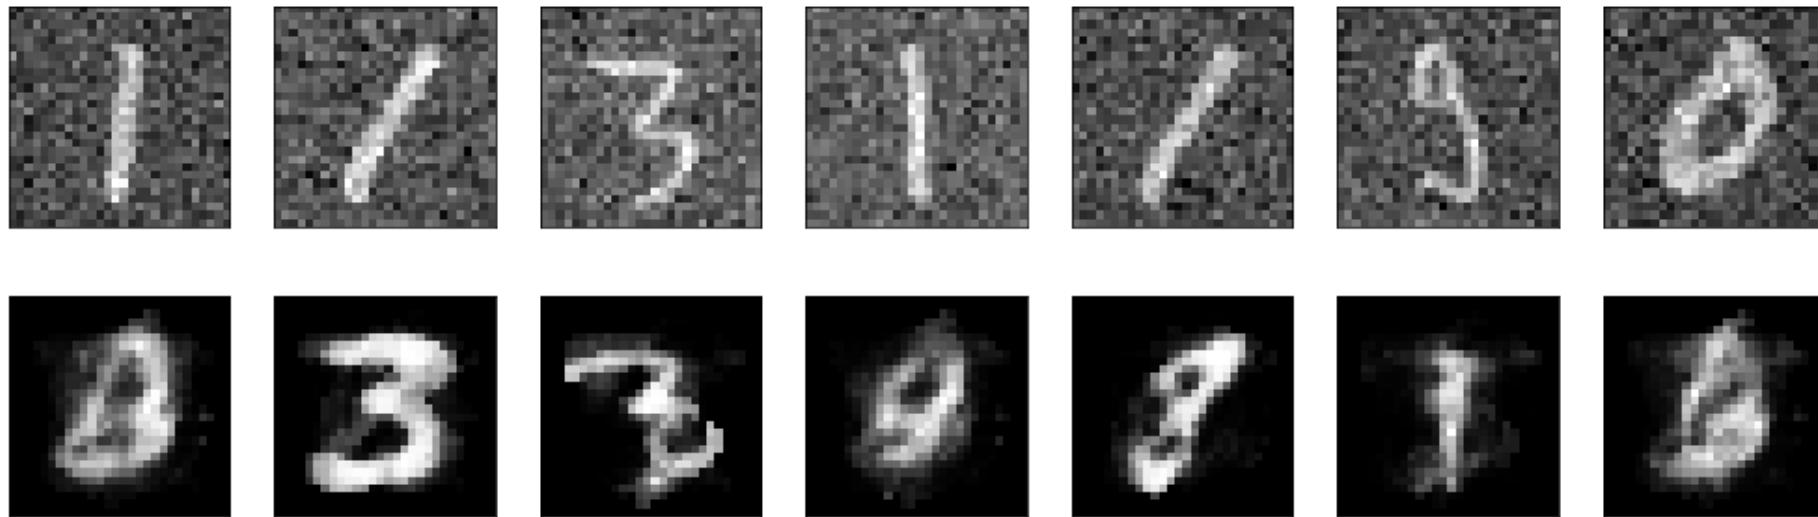


# DCGAN Training



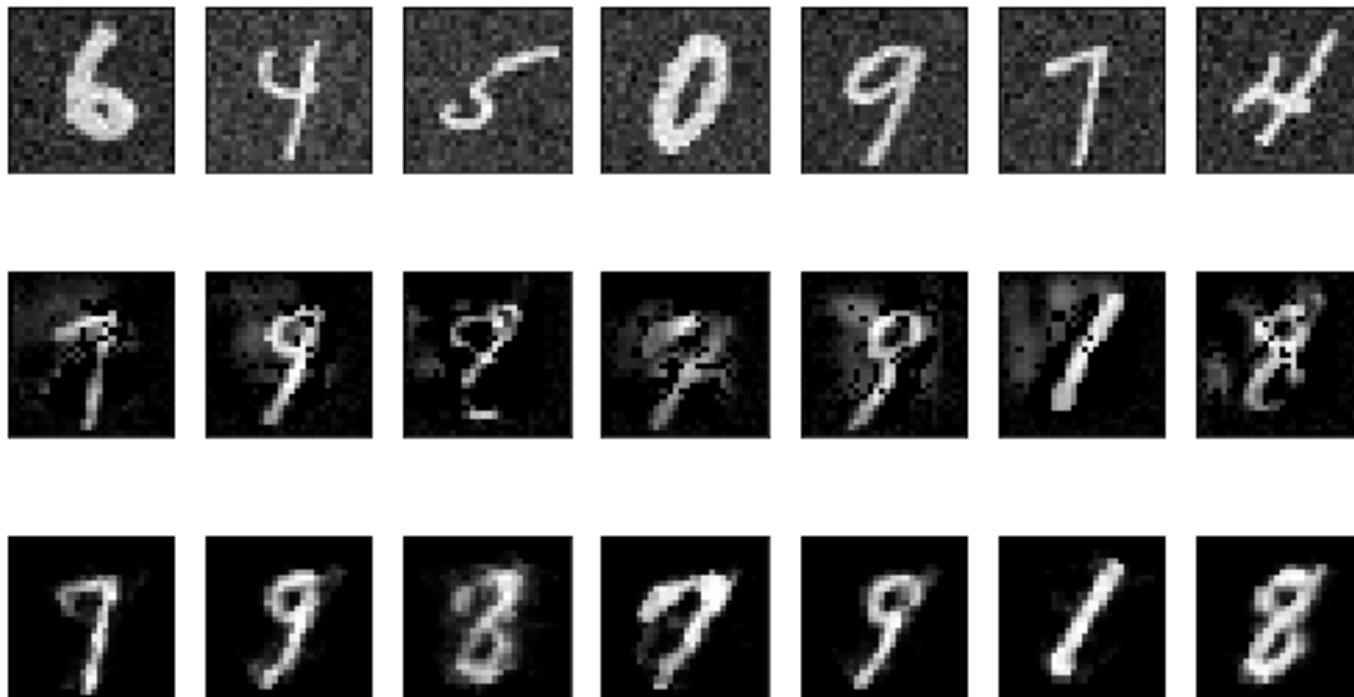
## DCGAN Output

```
Epoch 1750: Generator Loss: 0.042980, Discriminator Loss: 0.030590
Epoch 1800: Generator Loss: 0.042635, Discriminator Loss: 0.029989
Epoch 1850: Generator Loss: 0.039975, Discriminator Loss: 0.029853
Epoch 1900: Generator Loss: 0.038871, Discriminator Loss: 0.029369
Epoch 1950: Generator Loss: 0.037111, Discriminator Loss: 0.029259
Epoch 2000: Generator Loss: 0.035706, Discriminator Loss: 0.029093
```

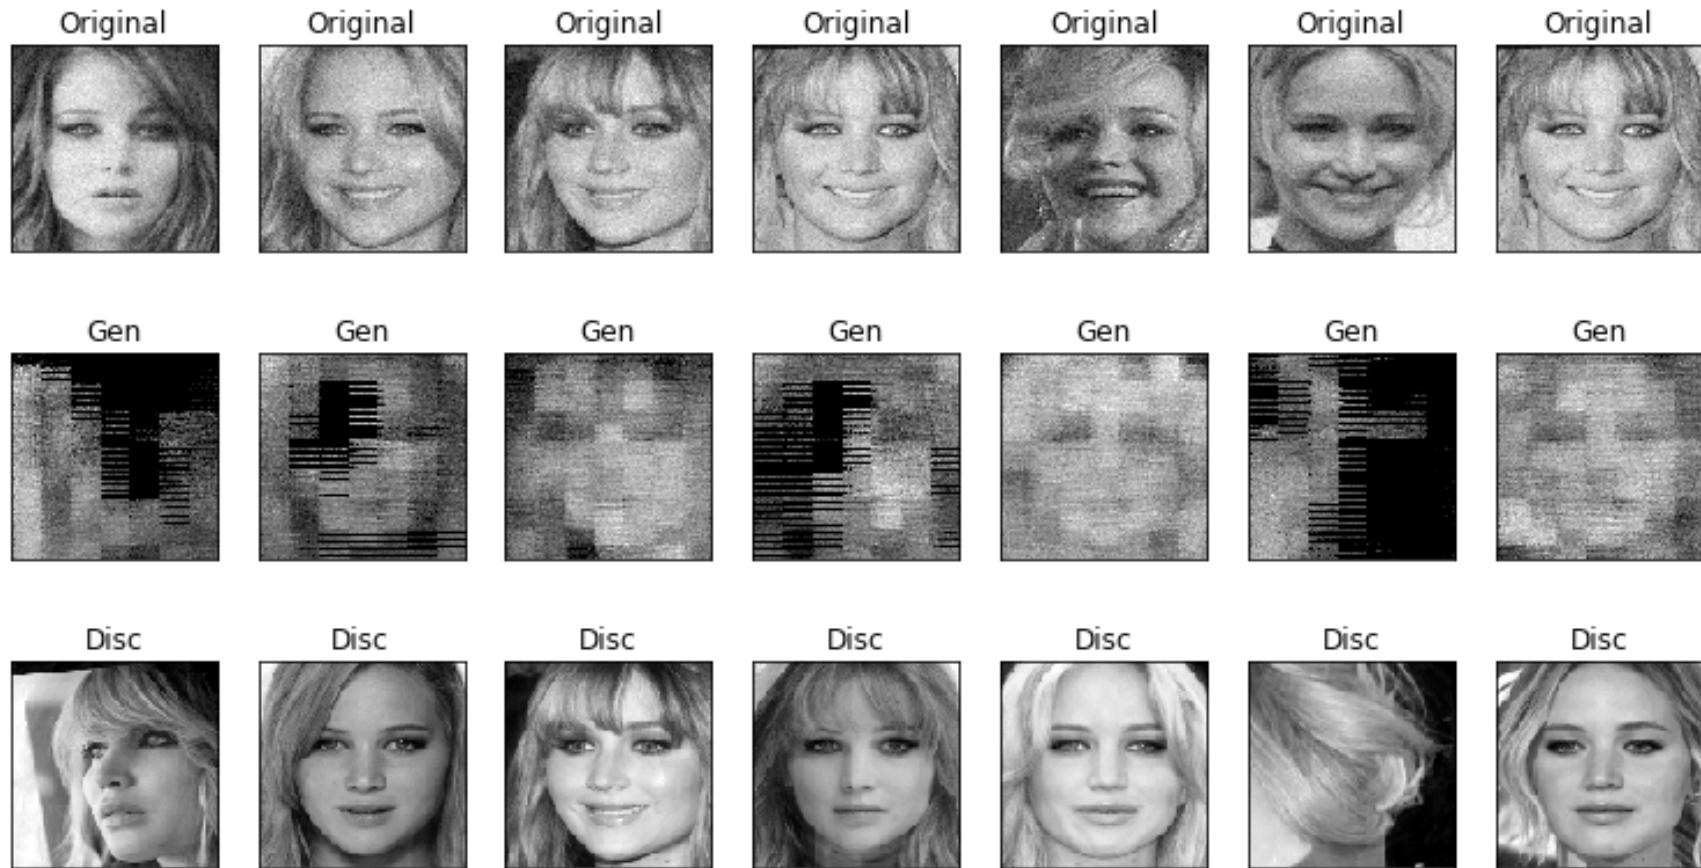


# DCGAN Output

Epoch 2850: Generator Loss: 0.051090, Discriminator Loss: 0.014070  
Epoch 2900: Generator Loss: 0.050746, Discriminator Loss: 0.013686  
Epoch 2950: Generator Loss: 0.050411, Discriminator Loss: 0.013317  
Epoch 3000: Generator Loss: 0.050048, Discriminator Loss: 0.012959



# DCGAN - Face Recognition

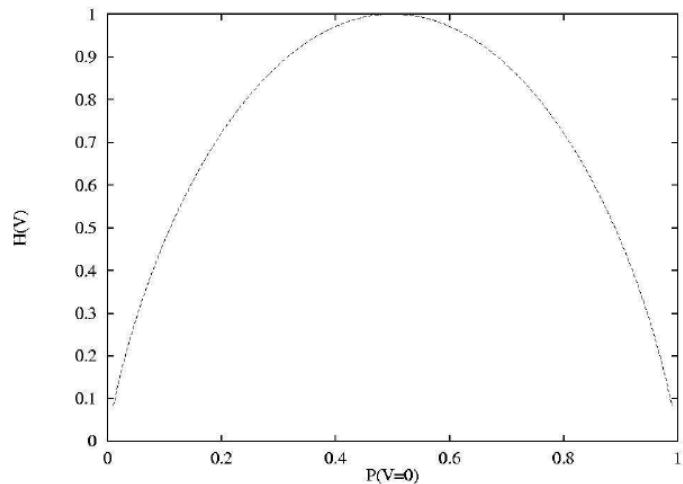


## Info GAN

- Based in Mutual Information and Entropy (uncertainty)

$$H(V) = \sum_{v=0}^1 -P(H = v) \lg P(H = v).$$

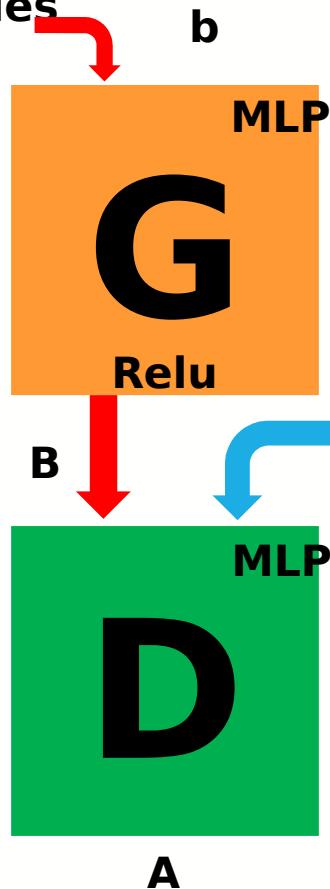
**Coin toss**



$$I(A;B) = H(B) - \sum_b P(B = b) \cdot H(A|B = b)$$

Latent Variables  
(add noise)

Fake



$$b = [1, 1, 1, 1, 1, 1, 1, 1, 1, 1] \quad N=10$$

$$H(B) = -\text{SumP}(B=b) \cdot \log P(B=b)$$

$$B = [0, 0, 0, 0, 0, 1, 1, 1, 1, 1]$$

$$\mathbf{P(B=b)} = 0.5$$

$$H(B) = -0.5 \cdot \log(0.5)$$

$$\mathbf{H(B)=-0.15}$$

Training Data  
Real

$$B = [0, 0, 0, 0, 0, 1, 1, 1, 1, 1]$$

$$A = [0, 0, 0, 0, 0, 1, 1, 1, 1]$$

$$P(A|B=b) = 4/5 = 0.8$$

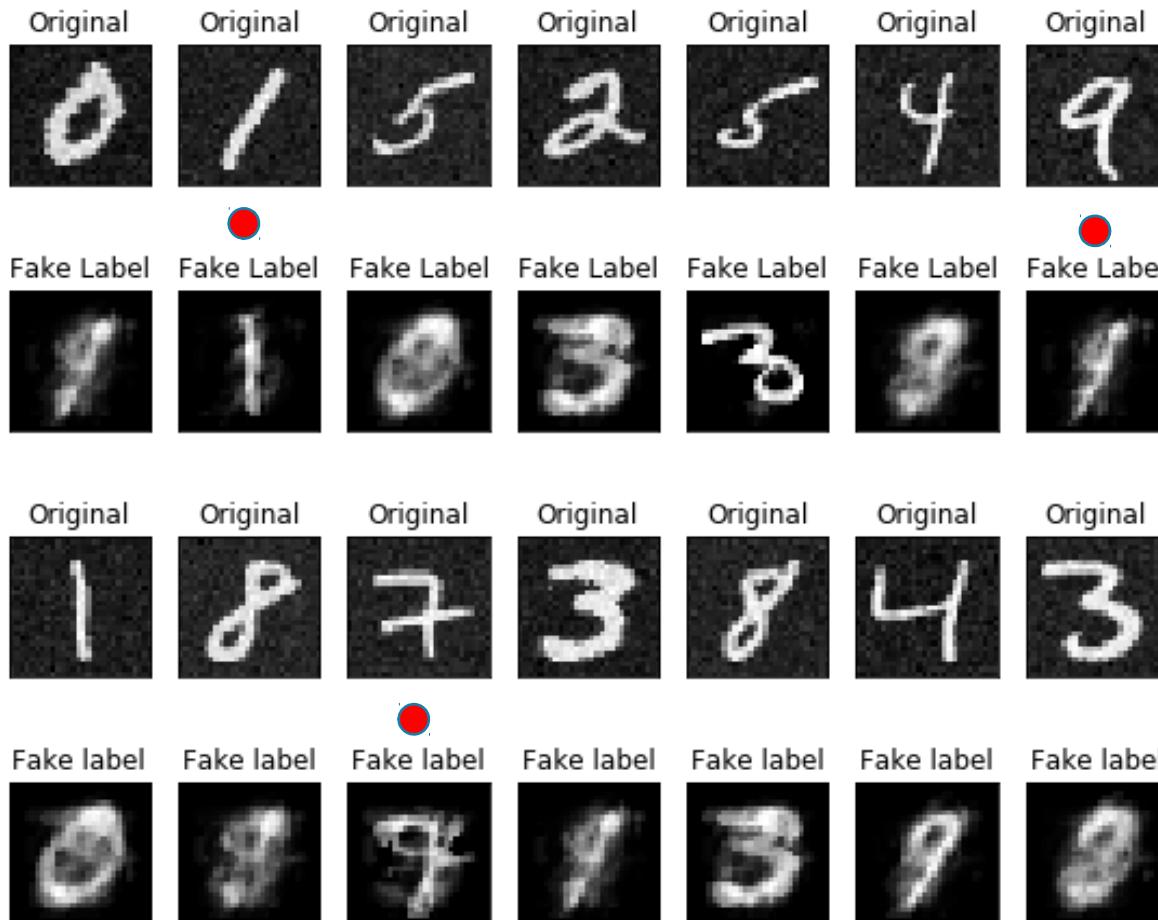
(subsample)

$$H(A|B=b) = -0.8 \cdot \log(0.8)$$

$$\mathbf{H(A|B=b)=-0.04}$$

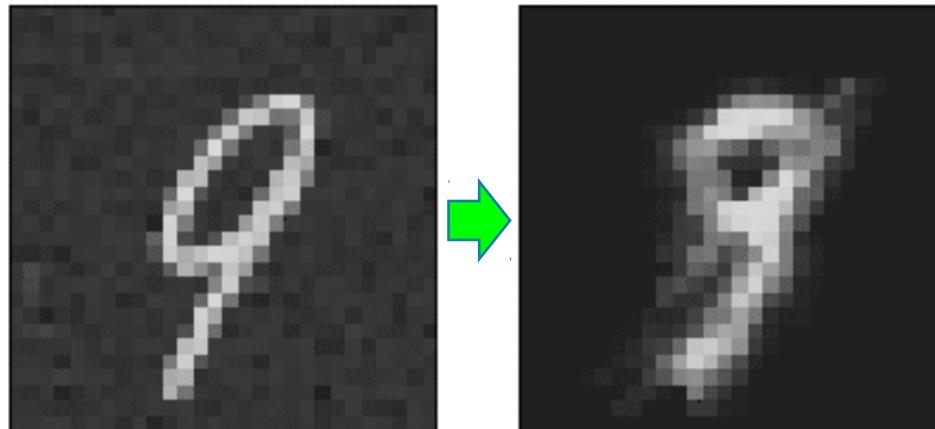
$$I(A; B) = H(B) - \sum_b P(B = b) \cdot H(A|B = b)$$

# Info GAN Output



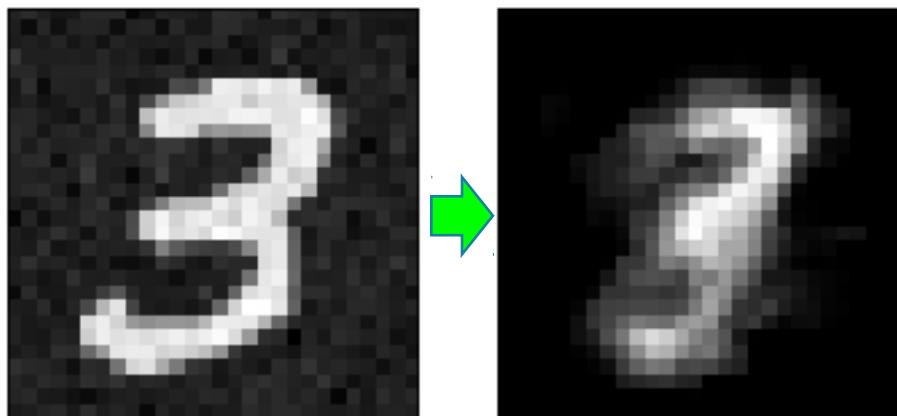
## Info GAN Output

```
In [135]: plt.figure(figsize=(6, 6))
...: ax = plt.subplot(1, 2, 1)
...: plt.imshow(x_train_noisy[0].reshape(28, 28))
...: ax.get_xaxis().set_visible(False)
...: ax.get_yaxis().set_visible(False)
...: ax = plt.subplot(1, 2, 2)
...: plt.imshow(np.array(g).reshape(60,28, 28)[0])
...: ax.get_xaxis().set_visible(False)
...: ax.get_yaxis().set_visible(False)
...: plt.show()
```

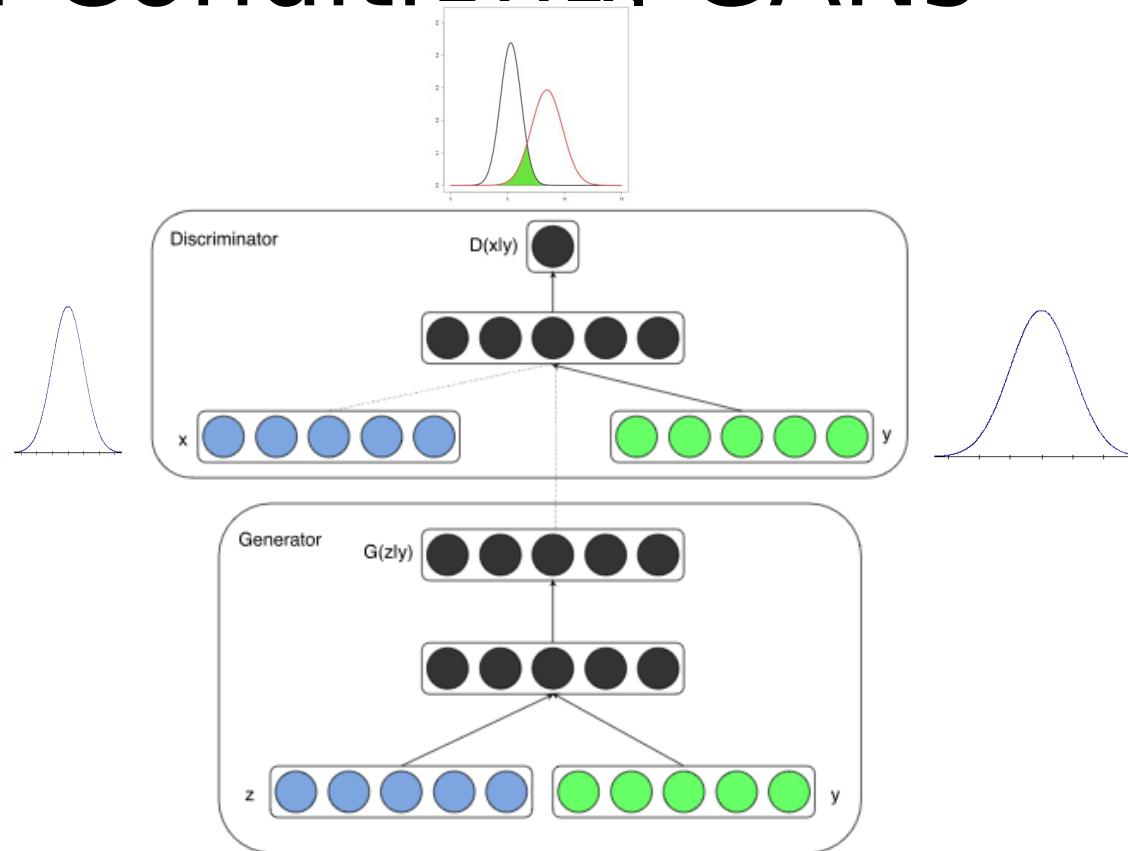


# Info GAN Output

```
In [186]: h=9
....: plt.figure(figsize=(6, 6))
....: ax = plt.subplot(1, 2, 1)
....: plt.imshow(x_train_noisy[h].reshape(28, 28))
....: ax.get_xaxis().set_visible(False)
....: ax.get_yaxis().set_visible(False)
....: ax = plt.subplot(1, 2, 2)
....: plt.imshow(np.array(g).reshape(60,28, 28)[h])
....: ax.get_xaxis().set_visible(False)
....: ax.get_yaxis().set_visible(False)
....: plt.show()
```



# cGAN: Conditional GANs

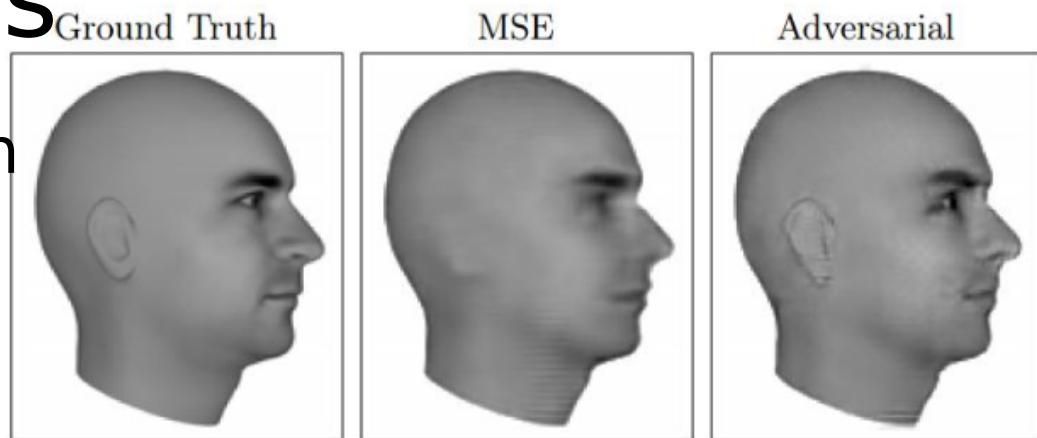


$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x|y)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z|y)))]$$

## GAN Examples

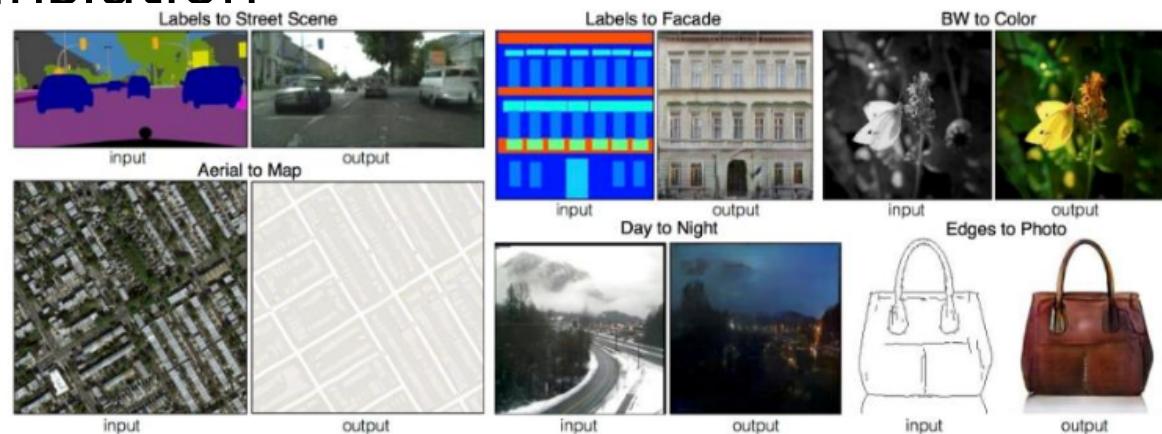
- Next frame prediction

Lotter et al, 2016



- Image-to-image translation

Isola et al, 2016



# GAN Examples

- Text-to-Image generation  
Inpainting

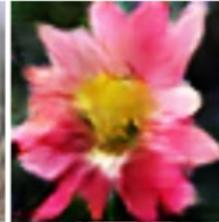
This bird is white with some black on its head and wings, and has a long orange beak



This bird has a yellow belly and tarsus, grey back, wings, and brown throat, nape with a black face



This flower has overlapping pink pointed petals surrounding a ring of short yellow filaments



(a) StackGAN  
Stage-I  
64x64  
images



(b) StackGAN  
Stage-II  
256x256  
images



Reed et al, 2016

## Image



(a) Input context

(b) Human artist



Pathak et al, 2016

# Cycle GAN: Style Transfer

 [junyanz / CycleGAN](#)



Monet ↪ Photos



Monet → photo

Zebras ↪ Horses



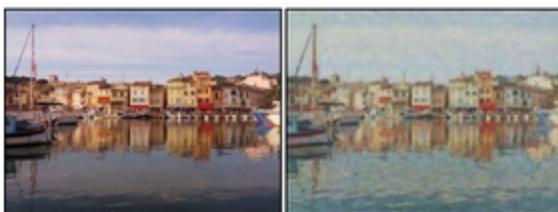
zebra → horse

Summer ↪ Winter

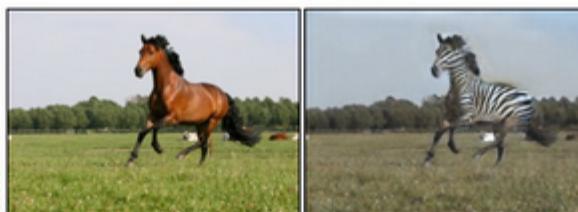


summer → winter

photo → Monet



horse → zebra



winter → summer





<https://github.com/RubensZimbres>

 Keras



# Workshop

---

Try different strategies in GANs